



SPEECH/VISION COGNITIVE SYSTEMS

Dr TIAN Jing

tianjing@nus.edu.sg



Module objective

Knowledge and understanding

- Understand the fundamentals of speech recognition systems, including statistical acoustic modelling, and end-to-end system using machine learning
- Understand basic concepts of vision cognitive systems

Key skills

- Design, build, implement and evaluate speech recognition approach in Python
- Design, build, implement and evaluate various visual question and answering approach in Python



Major reference

- [Introduction] CS131: ***Computer Vision: Foundations and Applications***,
http://vision.stanford.edu/teaching/cs131_fall1718/syllabus.html
- [Comprehensive] ***Computer Vision Crash Course***,
<https://filebox.ece.vt.edu/~jbhuang/>
- [Introduction] ***Automatic Speech Recognition***,
https://github.com/ekapolc/ASR_course
- [Comprehensive] CS224S, ***Spoken Language Processing***,
<http://web.stanford.edu/class/cs224s/>



Topics

- Vision cognition systems
- Speech recognition systems
- Workshop: Design and build speech recognition system in Python

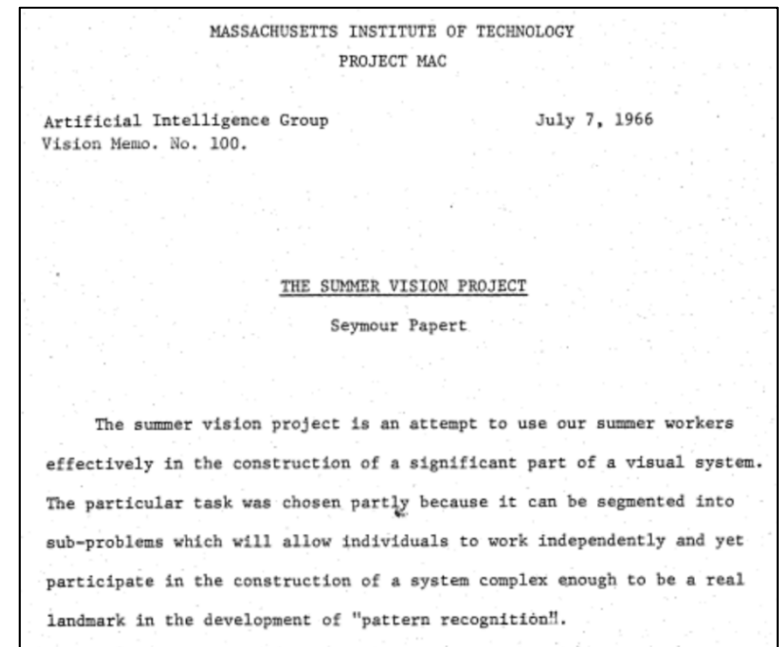
Vision cognition

The first computer vision project in 1966.

Abstract: The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of “pattern recognition”.

Tasks

- **Figure ground:** Divide a picture into regions such as likely objects, likely background areas.
- **Region description:** Analysis of shape and surface properties.
- **Object identification:** Name objects by matching them with a vocabulary of known objects.

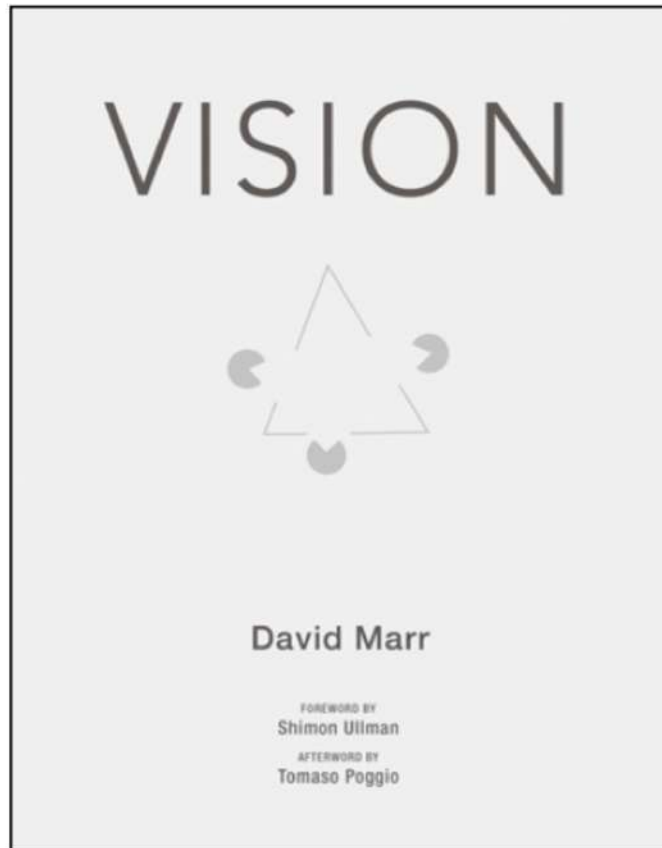


Reference: <http://people.csail.mit.edu/brooks/idocs/AIM-100.pdf>



Vision cognition: Marr theory

I am not sure that Marr would agree, but I am tempted to add learning as the very top level of understanding, above the computational level. (T. Poggio, 2010)



Learning

Computational

- Computations relating inputs to outputs

Algorithmic

- How the computation is executed at the level of information processing

Implementational

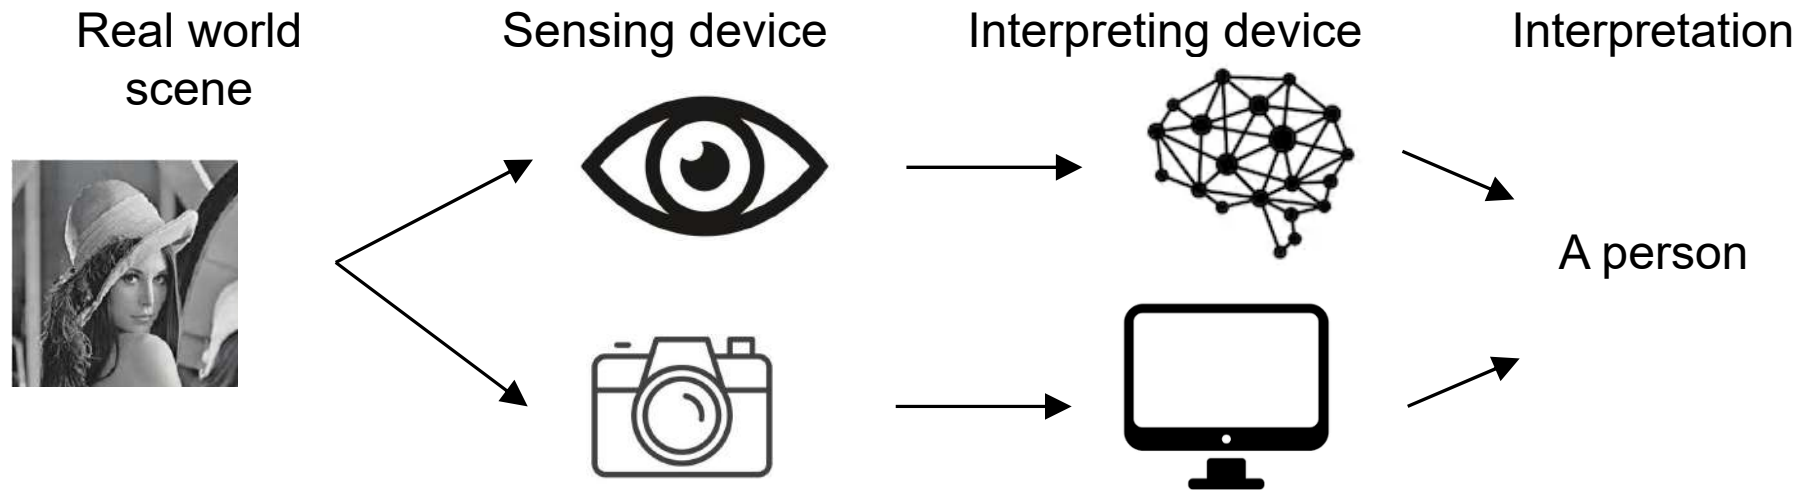
- How algorithm is embodied as a physical process

Reference

- [https://en.wikipedia.org/wiki/David_Marr_\(neuroscientist\)](https://en.wikipedia.org/wiki/David_Marr_(neuroscientist))
- T. Poggio, Vision (2010, The MIT Press), Afterword, P.367

Thinking humanly

- Humans use their eyes and brains to visually sense the world.
- Computers use their cameras and computation to visually sense the world.



Computers	Brains
Fixed architecture	Evolving architecture
Modular, (primarily) serial	Massively parallel
Separate hardware, software	No distinction between hardware and software
Separate computation, memory	No distinction between computation and memory

Reference: <http://scienceblogs.com/developingintelligence/2007/03/27/why-the-brain-is-not-like-a-co/>



Key cognitivist vision tasks

A concept is named entity, e.g., cat, human

- Learn concepts given labeled examples
- Localize concepts given labeled examples
- Count concepts
- Search for examples similar to this concept
- Explain evidence for concepts
- Estimate variance intrinsic or extrinsic to concept

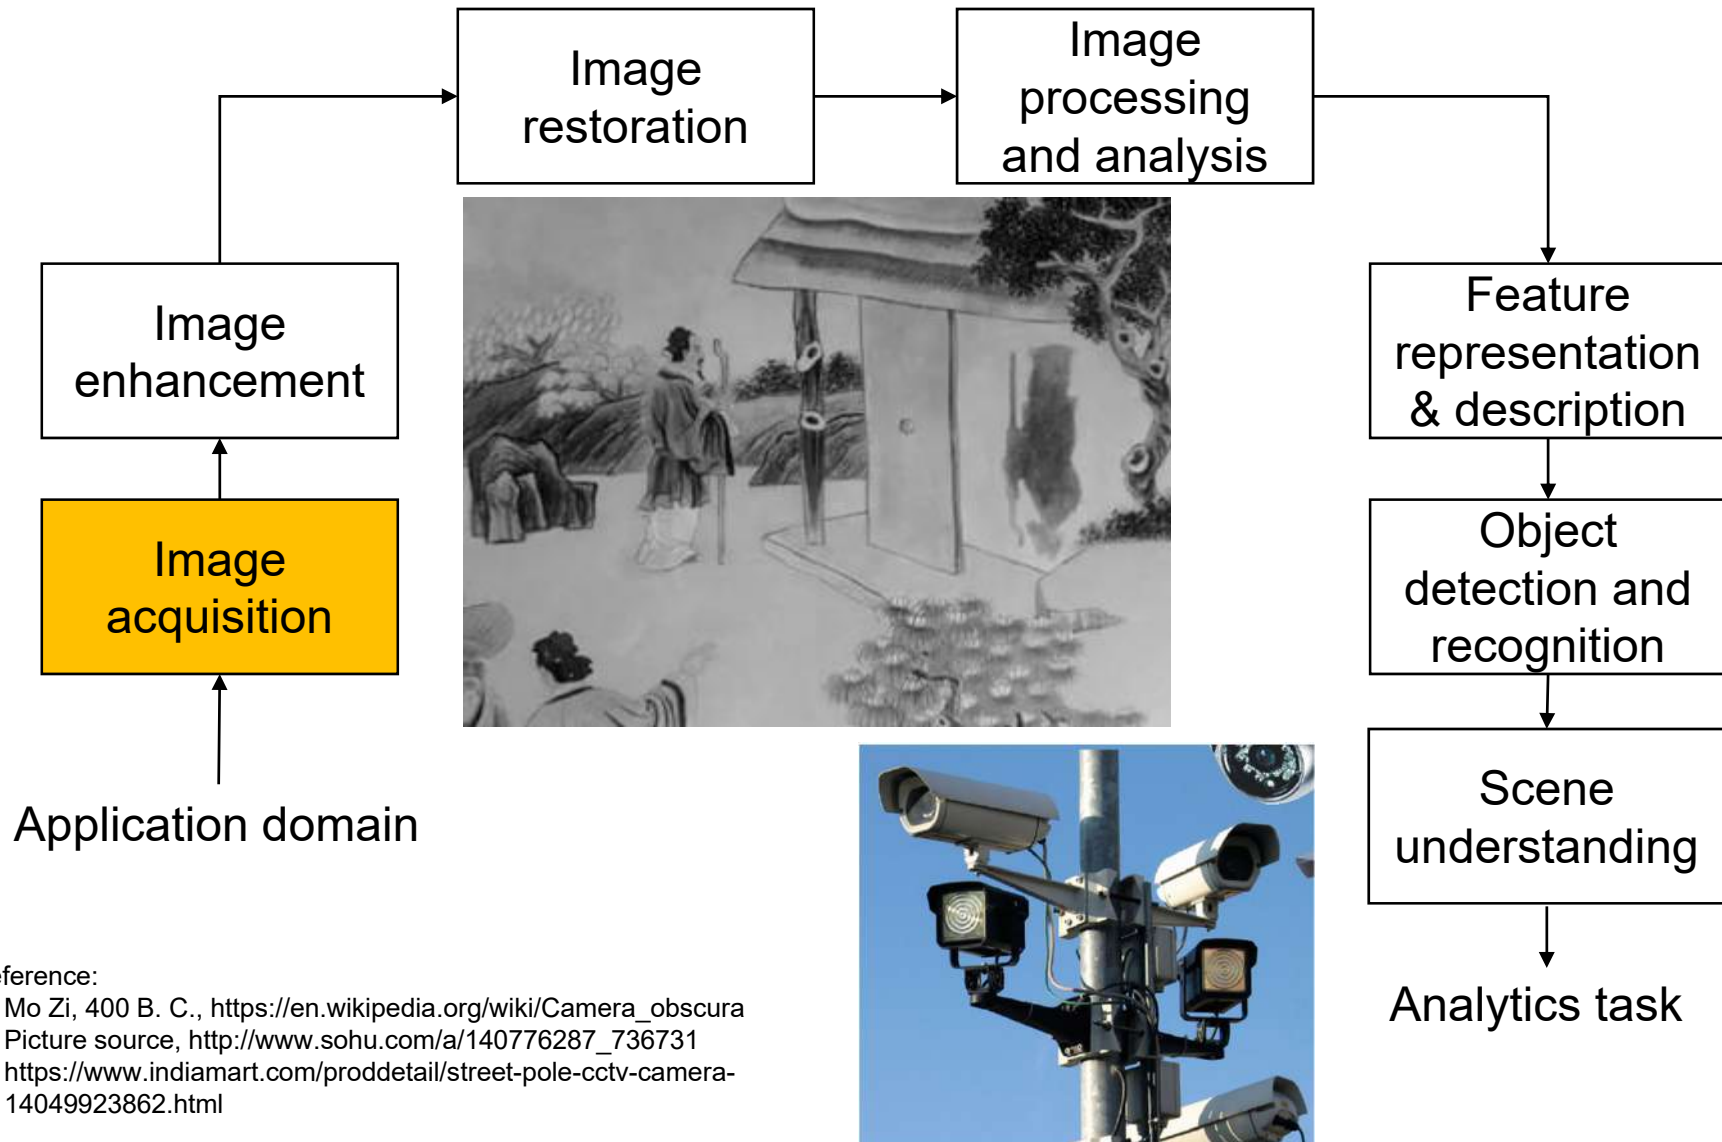


Photo: <https://cathumor.net/human-i-request-your-assistance/>

Behaviour decision	Cat should take off hanger
High-level interpretation	Cat hung on hanger
Scene description	Cat on the flat floor
Visual objects	Car, hanger
Integrated features	Histogram of color/gradients
Low-level features	Colors and textures
Pre-processing	Enhance contrast of images



Vision cognitive system pipeline

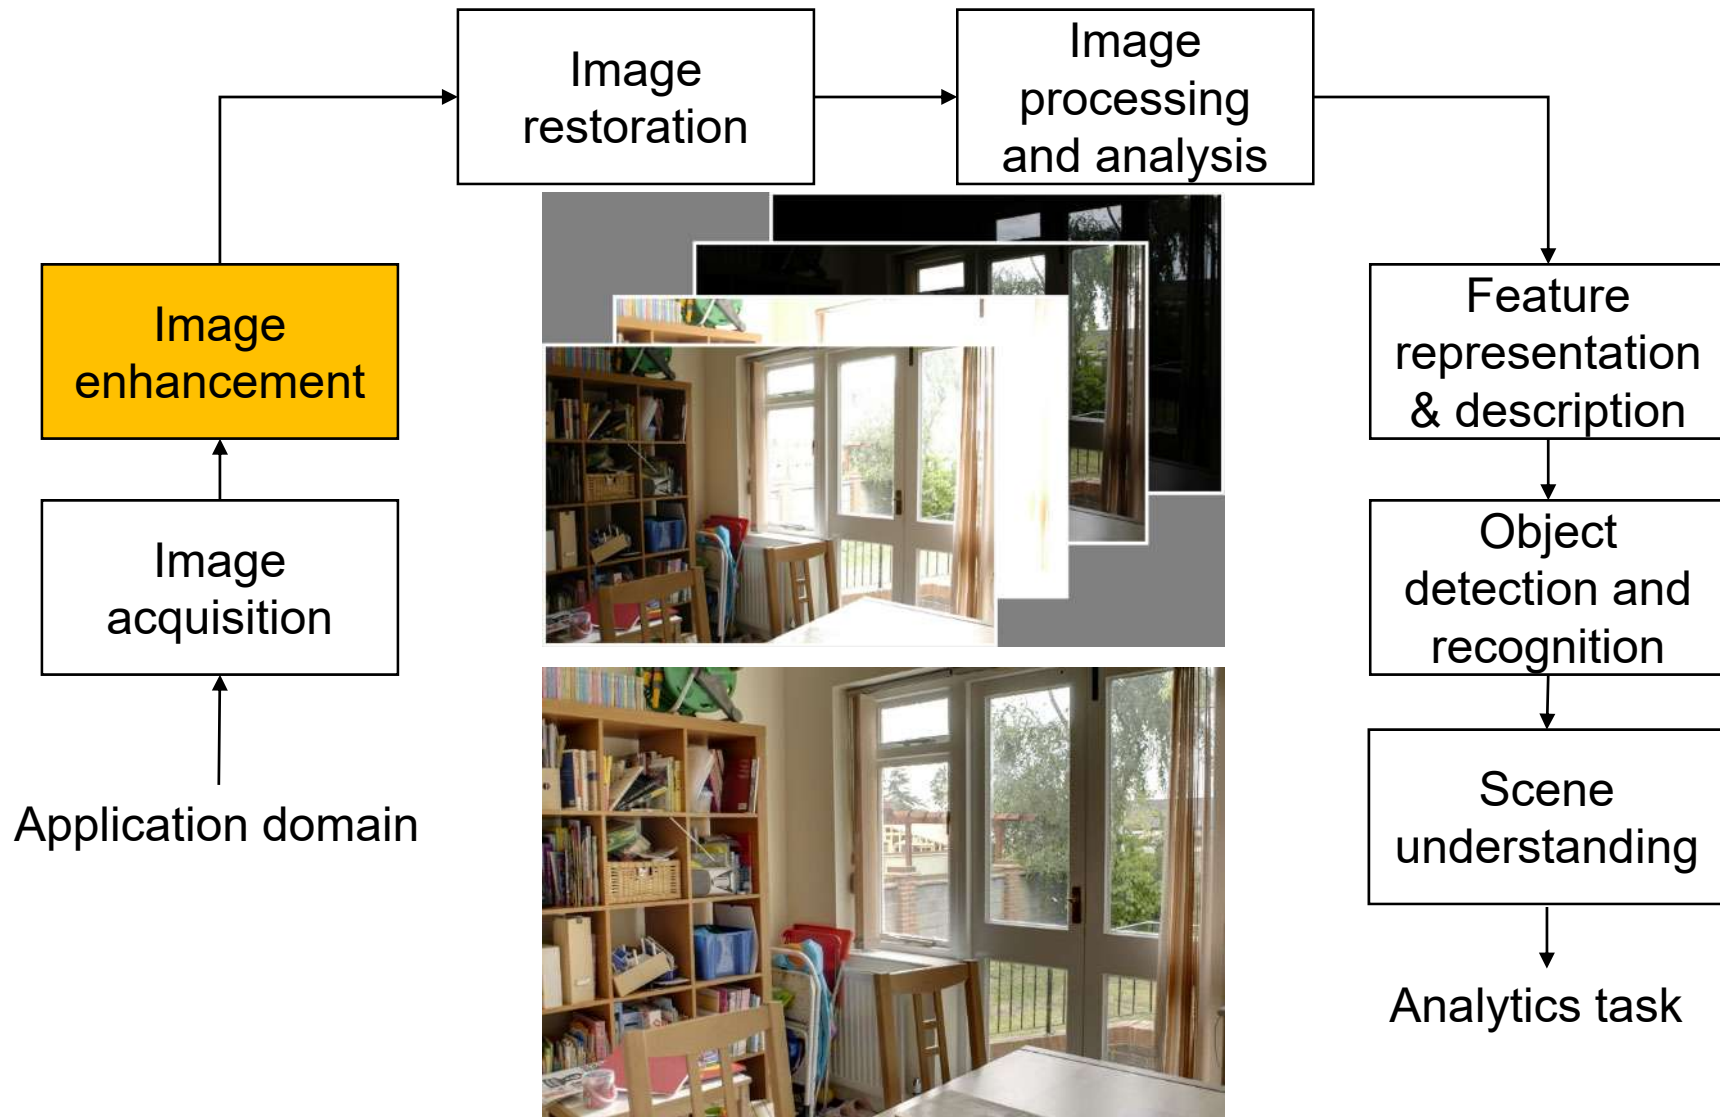


Reference:

- Mo Zi, 400 B. C., https://en.wikipedia.org/wiki/Camera_obscura
- Picture source, http://www.sohu.com/a/140776287_736731
- <https://www.indiamart.com/proddetail/street-pole-cctv-camera-14049923862.html>



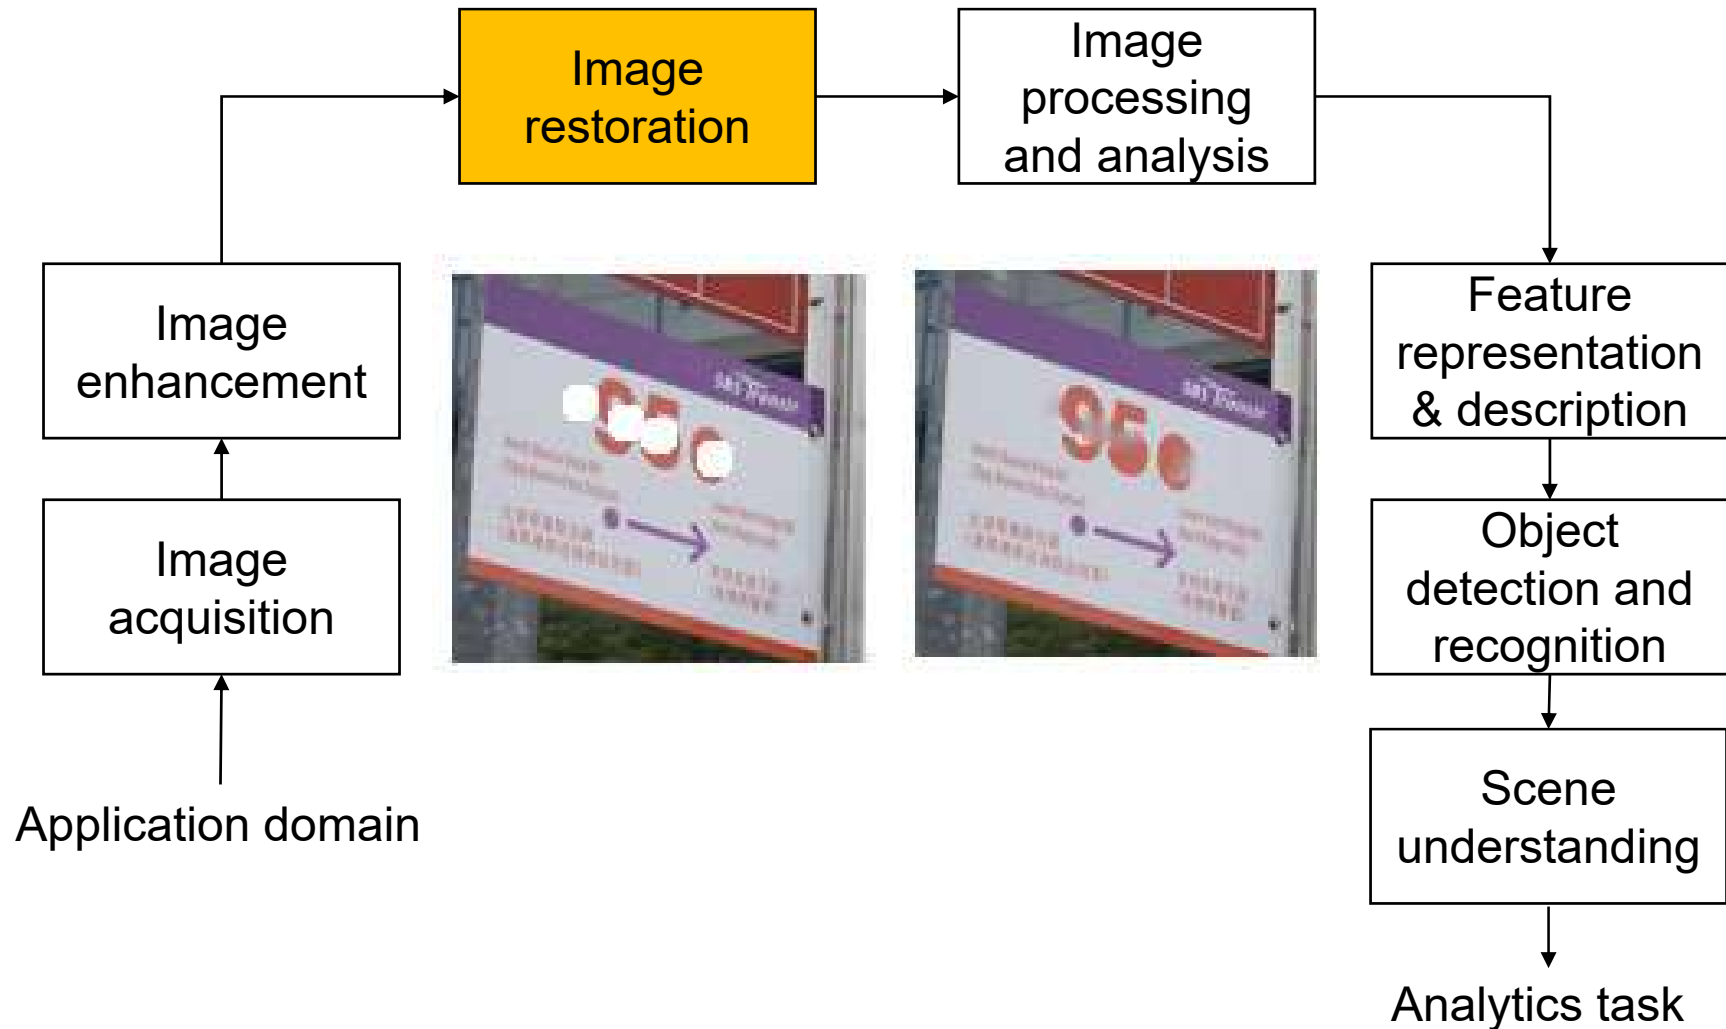
Vision cognitive system pipeline



Online demo: <http://ipolcore.ipol.im/demo/clientApp/demo.html?id=230>



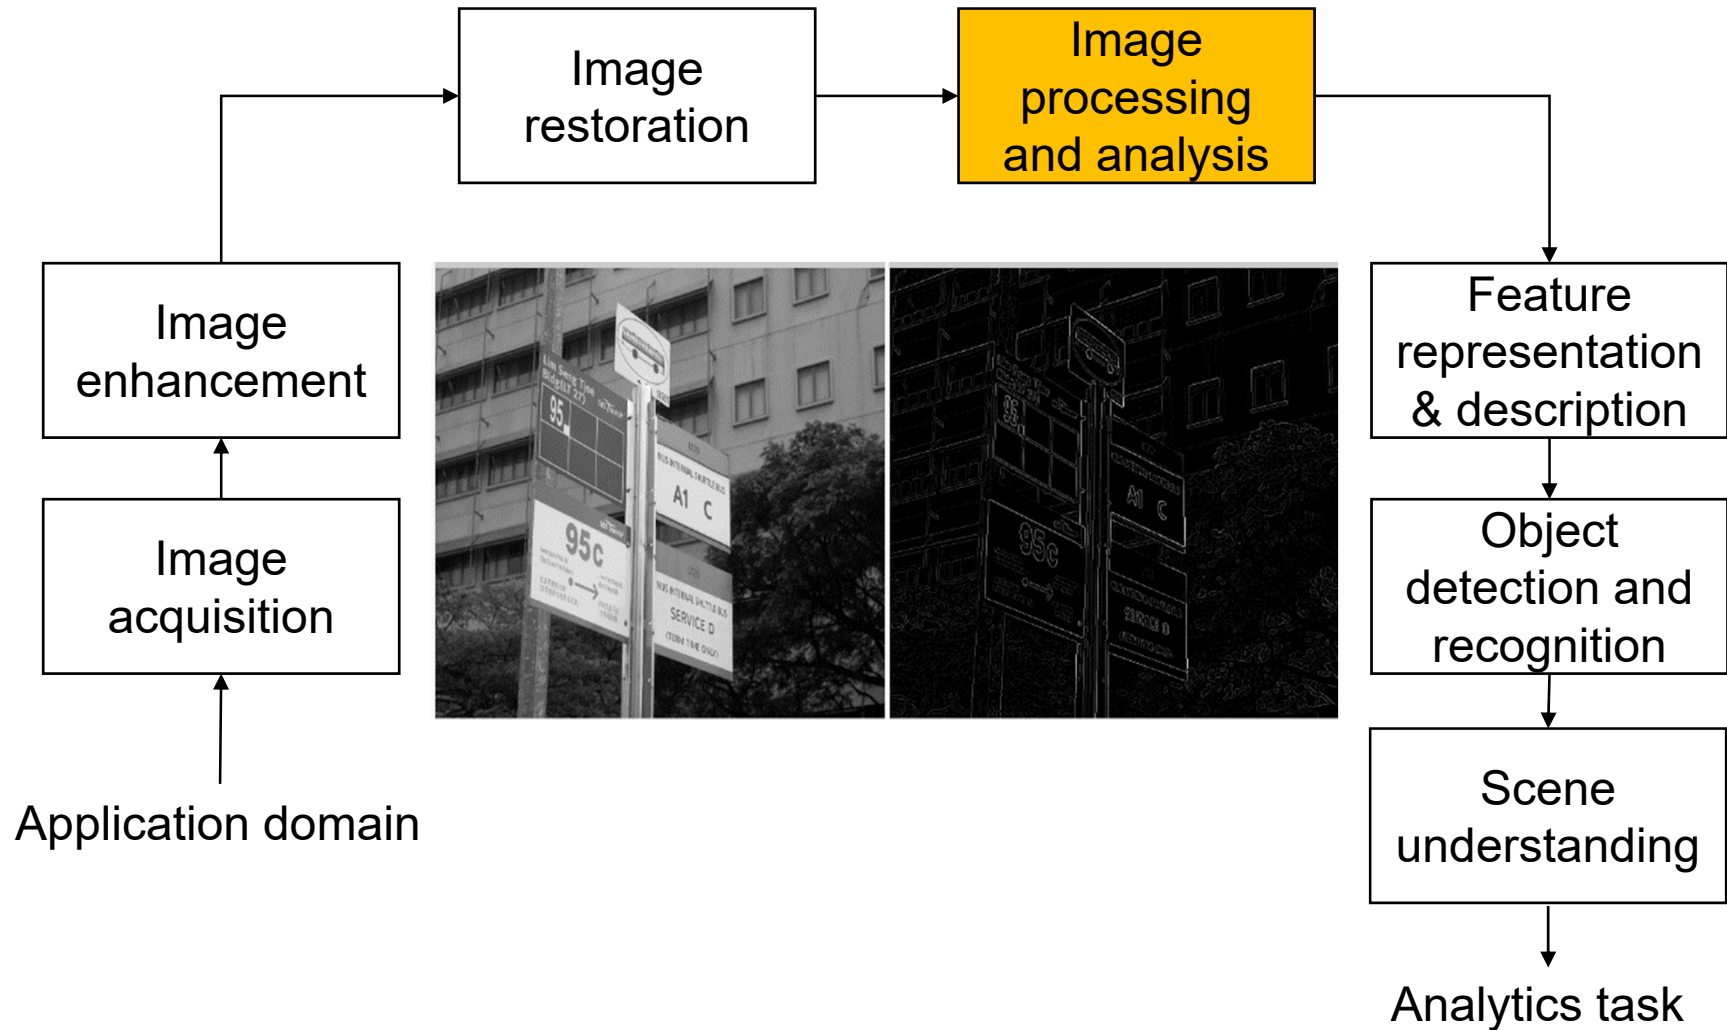
Vision cognitive system pipeline



Online demo: <http://demo.ipol.im/demo/54/>



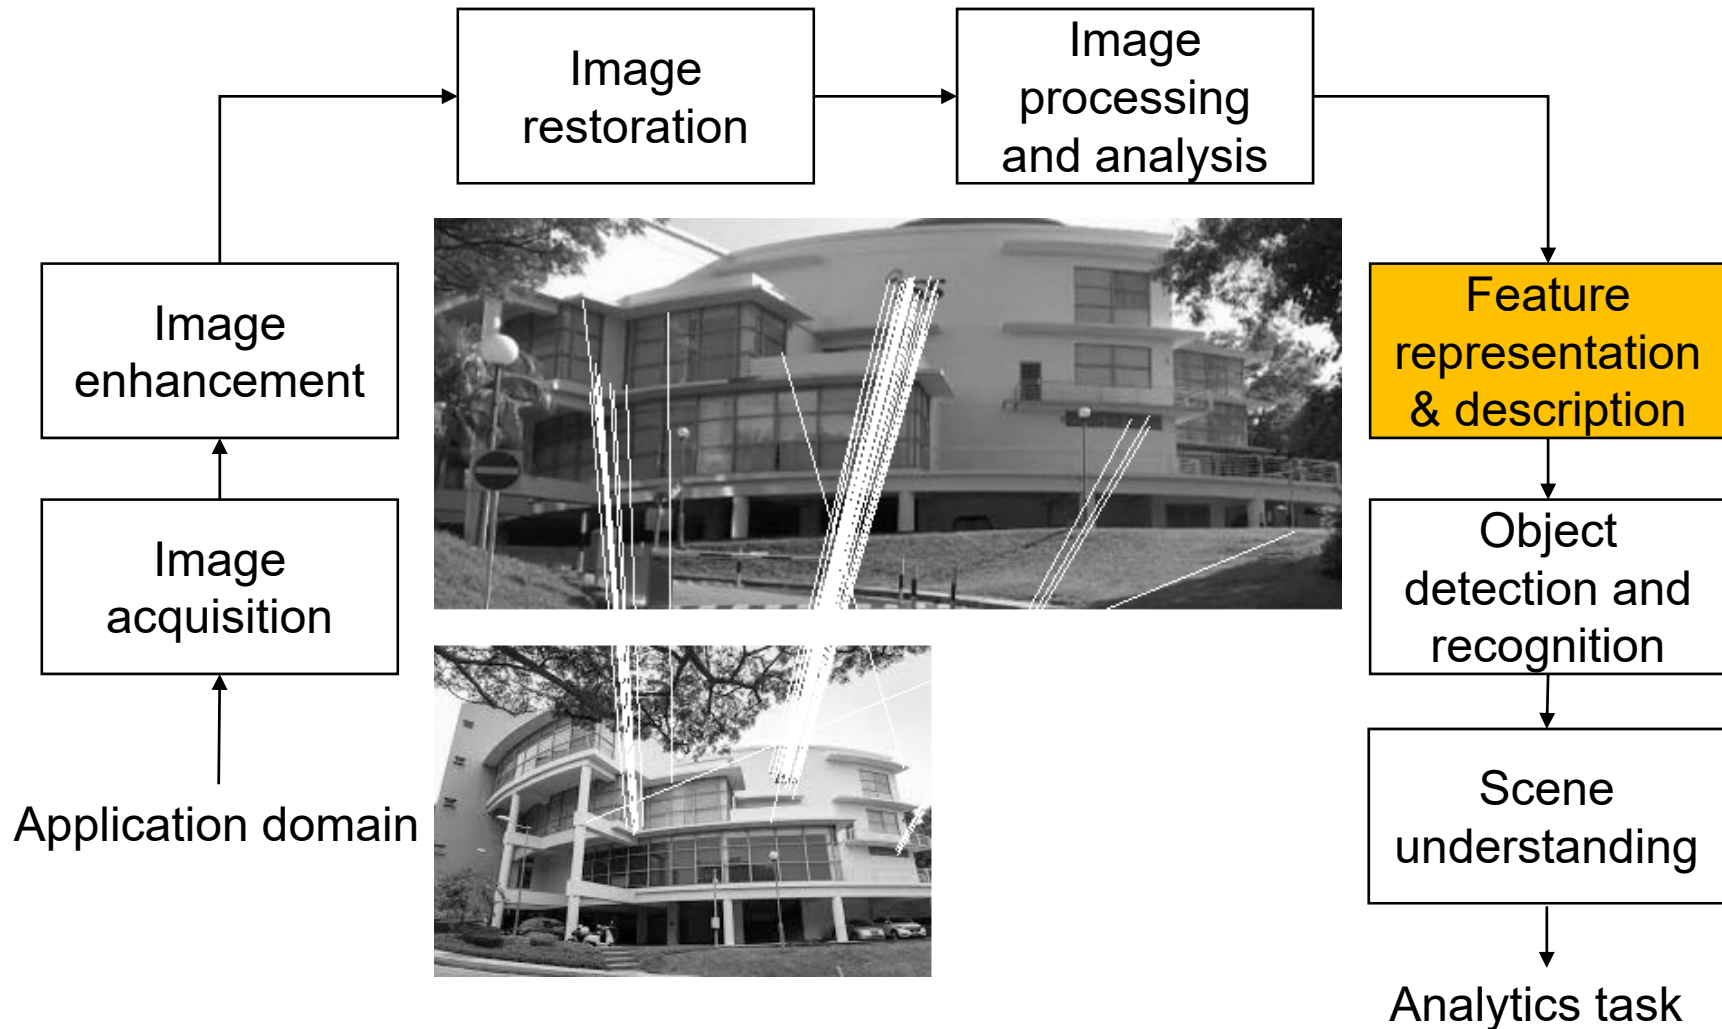
Vision cognitive system pipeline



Online demo: <http://bigwww.epfl.ch/demo/ip/demos/edgeDetector/>



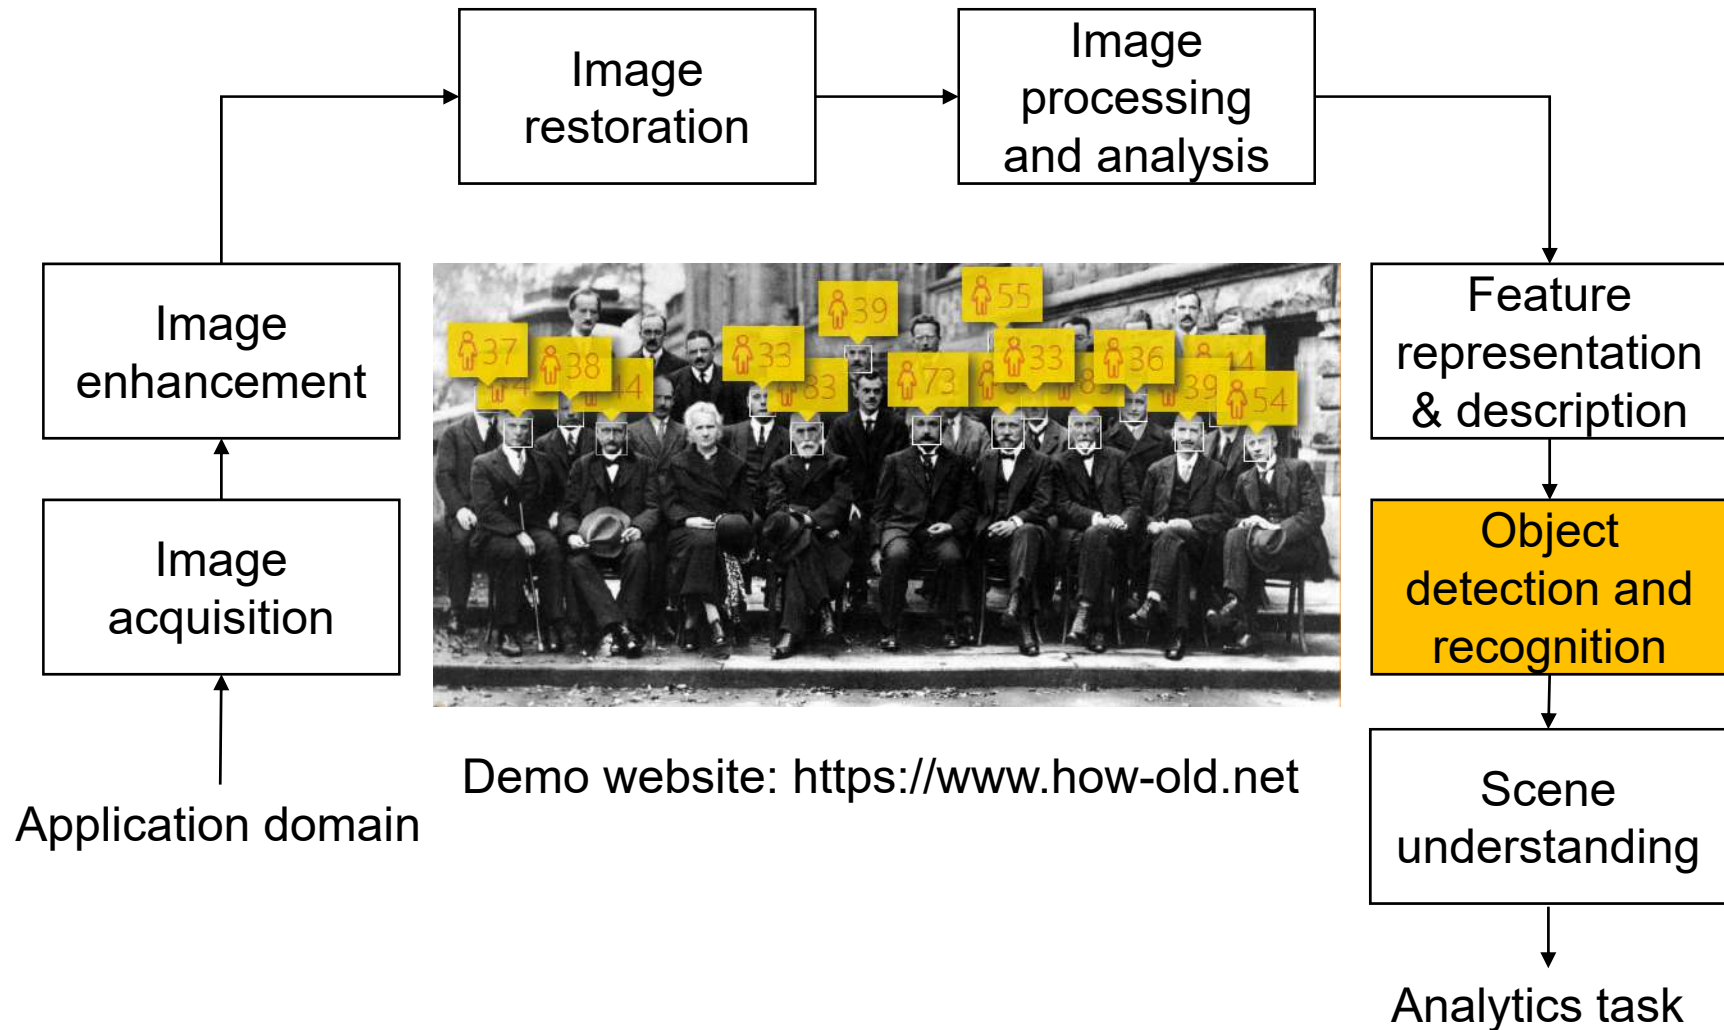
Vision cognitive system pipeline



Online demo: http://demo.ipol.im/demo/my_affine_sift/

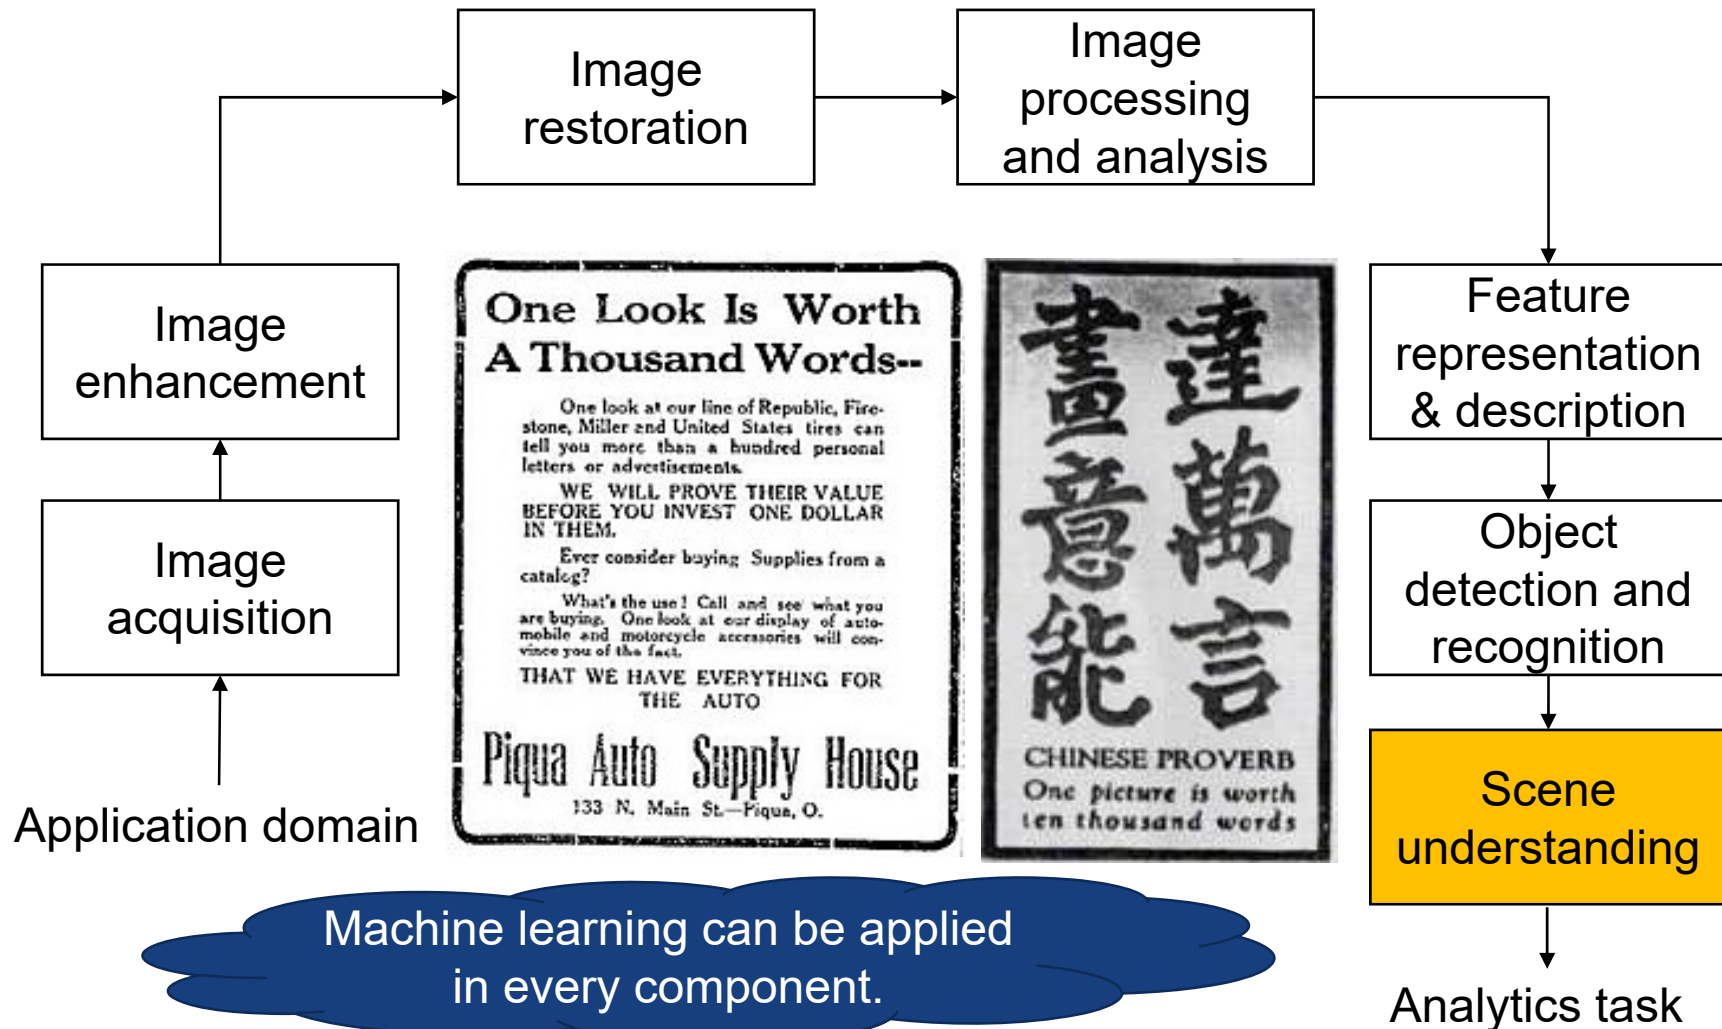


Vision cognitive system pipeline





Vision cognitive system pipeline



Reference

- https://en.wikipedia.org/wiki/A_picture_is_worth_a_thousand_words
- <https://www.phrases.org.uk/meanings/a-picture-is-worth-a-thousand-words.html>



Automatic speech recognition



Amazon Echo
2015



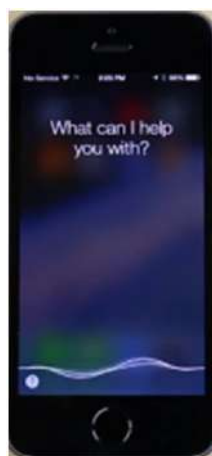
Google Home
2016



Facebook M
2015



Anki Cozmo
2016



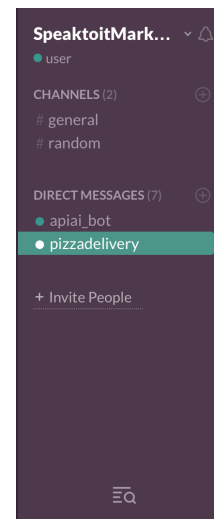
Apple
Siri
2011



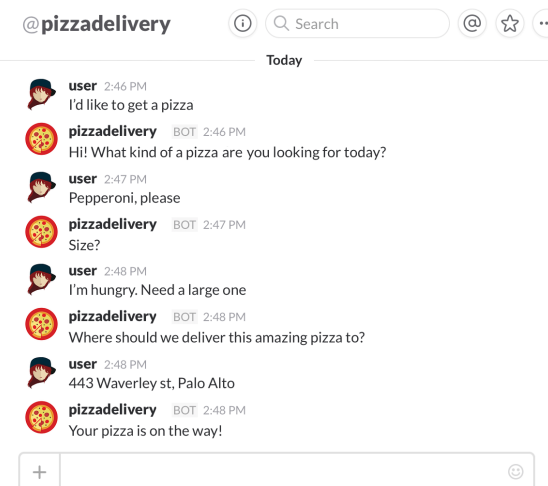
Google
Assistant
2016



Microsoft
Cortana
2014



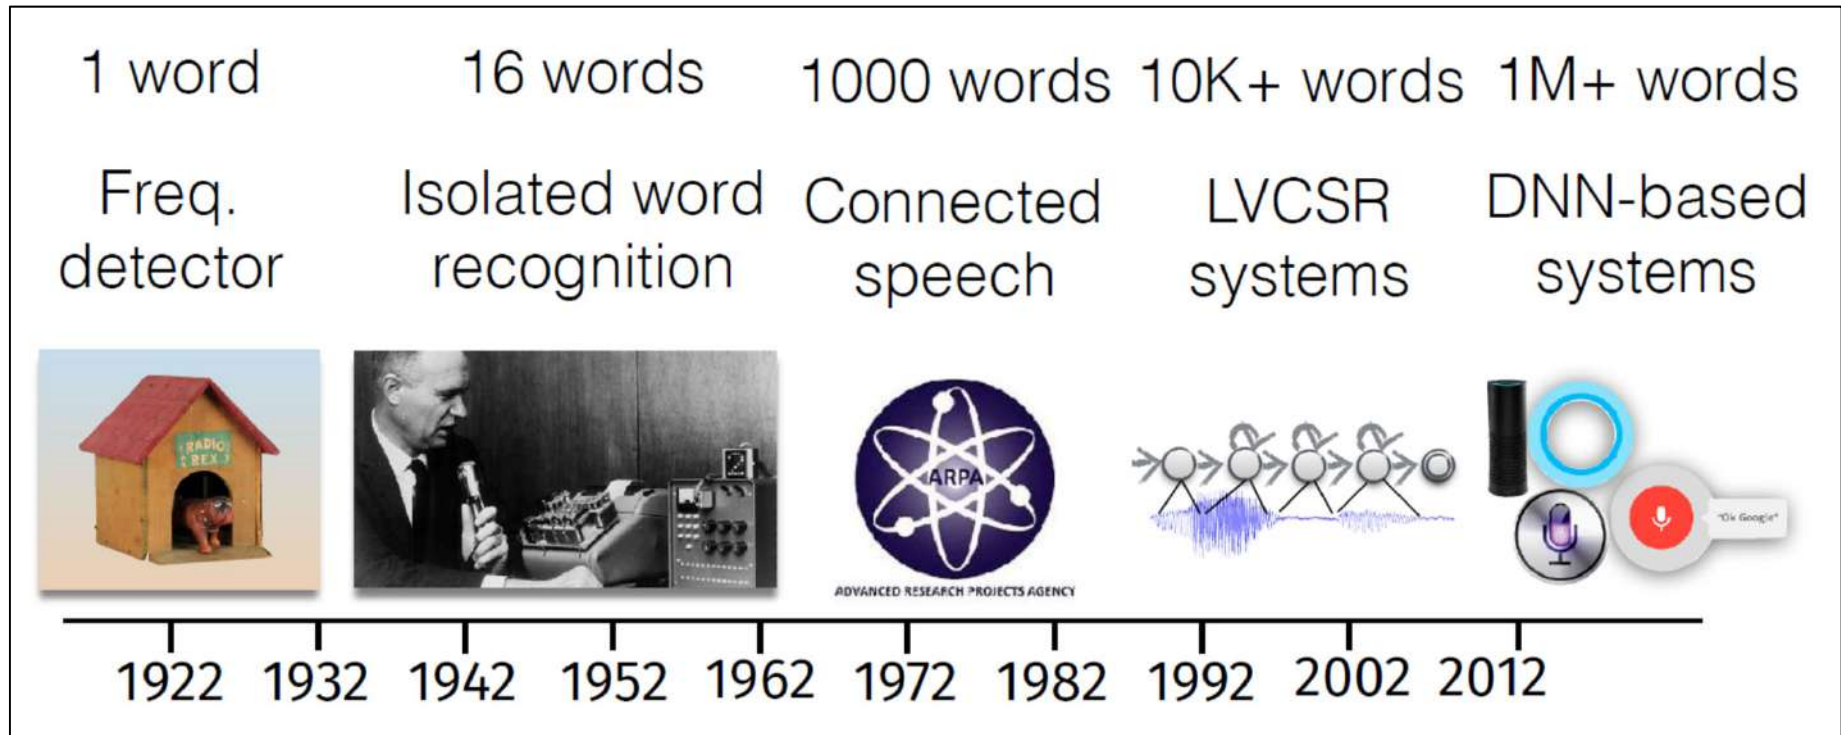
Slack Bot API
2015



Source: CS224S Spoken Language Processing, <http://web.stanford.edu/class/cs224s/>



Automatic speech recognition



More introductions on history of automatic speech recognition can be found at

- <https://ileriseviye.wordpress.com/2011/02/17/speech-recognition-in-1920s-radio-rex-the-first-speech-recognition-machine/>
- <https://machinelearning-blog.com/2018/09/07/a-brief-history-of-asr-automatic-speech-recognition/>

Source: Automatic Speech Recognition (CS753), Lecture 1: Introduction to Statistical Speech Recognition, <https://www.cse.iitb.ac.in/~pjyothi/cs753/>



Automatic speech recognition

Behind the Mic: The Science of Talking with Computers, (6 minutes)

<https://www.youtube.com/watch?v=yxxRAHVtafl>

Language is easy for humans to understand (most of the time), but not so easy for computers. This video talks about speech recognition, language understanding, neural nets, and using our voices to communicate with the technology around us.





Automatic speech recognition

- **Human-machine Interaction**
 - Automatic Speech Recognition
 - Speech Synthesis / Text-to-Speech (TTS)
 - Natural Language Understanding (NLU)
 - Natural Language Generation (NLG)
- **Telecommunication**
 - Speech Coding
- **Language**
 - Statistical Machine Translation (SMT)
- **Language Acquisition**
 - Pronunciation Training
- **Security/Forensics**
 - Speaker ID
 - Speaker Verification
- **Medical Applications**
 - Diagnosis of Diseases
- **Information Retrieval**
 - Video/Audio Transcribing
 - Audio/Text Summarizing
- **Speech Manipulation**
 - Speaking Rate Adjusting

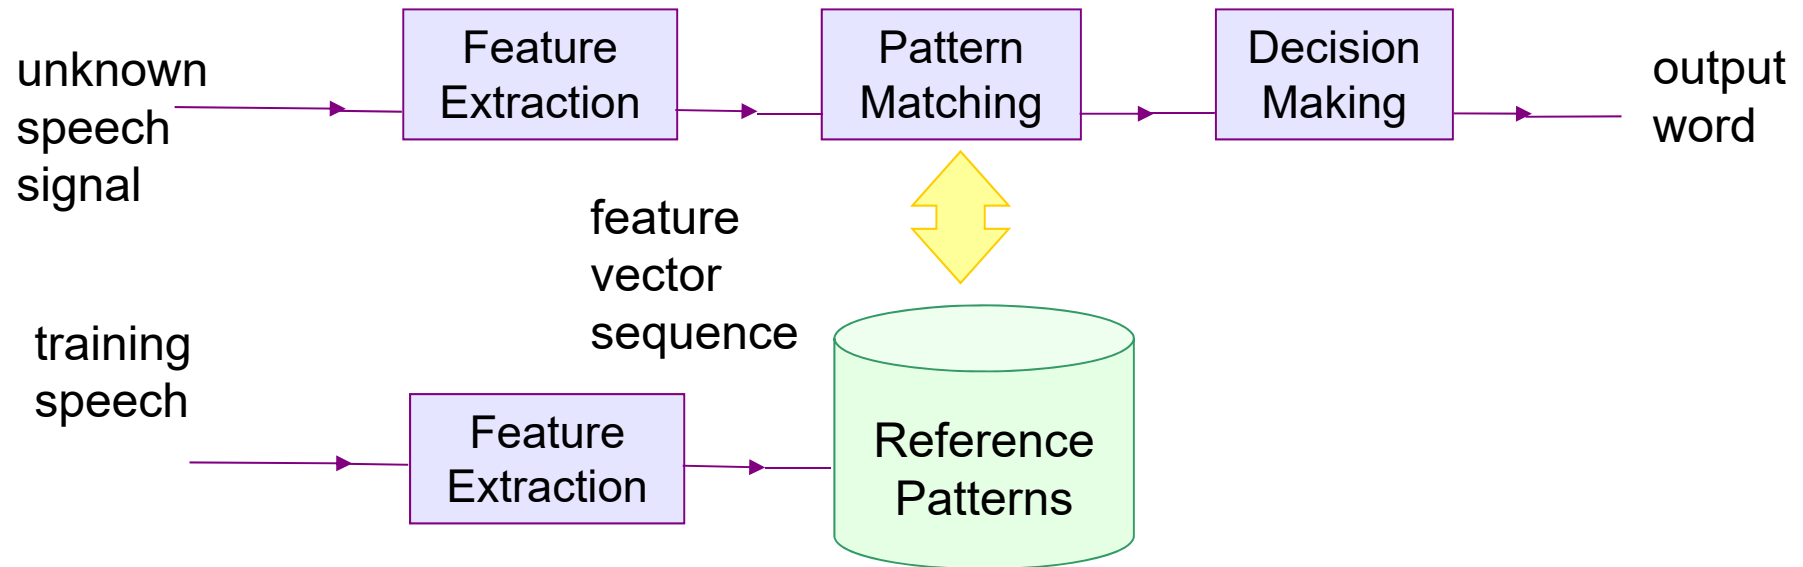


Automatic speech recognition

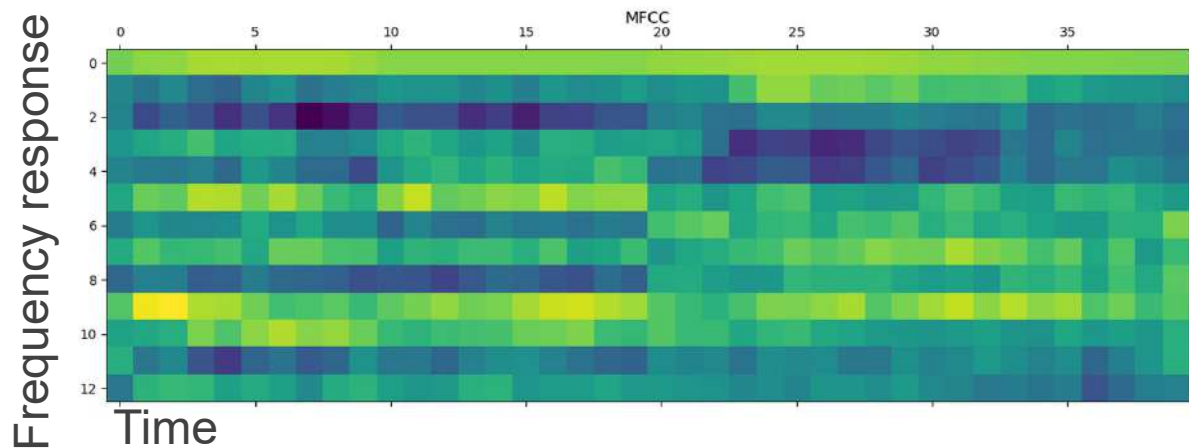
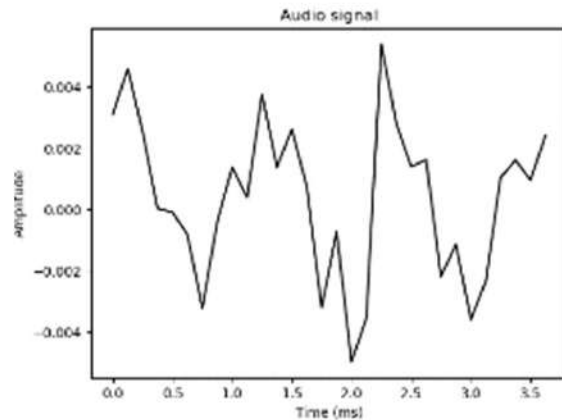
Challenges of speech recognition

- **Style:** Read speech or spontaneous (conversational) speech?
- **Continuous natural speech** or **command & control**?
- **Speaker characteristics:** Rate of speech, accent, prosody (stress, intonation), speaker age, pronunciation variability even when the same speaker speaks the same word
- **Channel characteristics:** Background noise, room acoustics, microphone properties, interfering speakers
- **Task specifics:** Vocabulary size (very large number of words to be recognized), language-specific complexity, resource limitations

Speech recognition as pattern recognition problem



Speech signal represented in time domain and frequency-domain features, e.g., *Mel frequency cepstral coefficient (MFCC)*

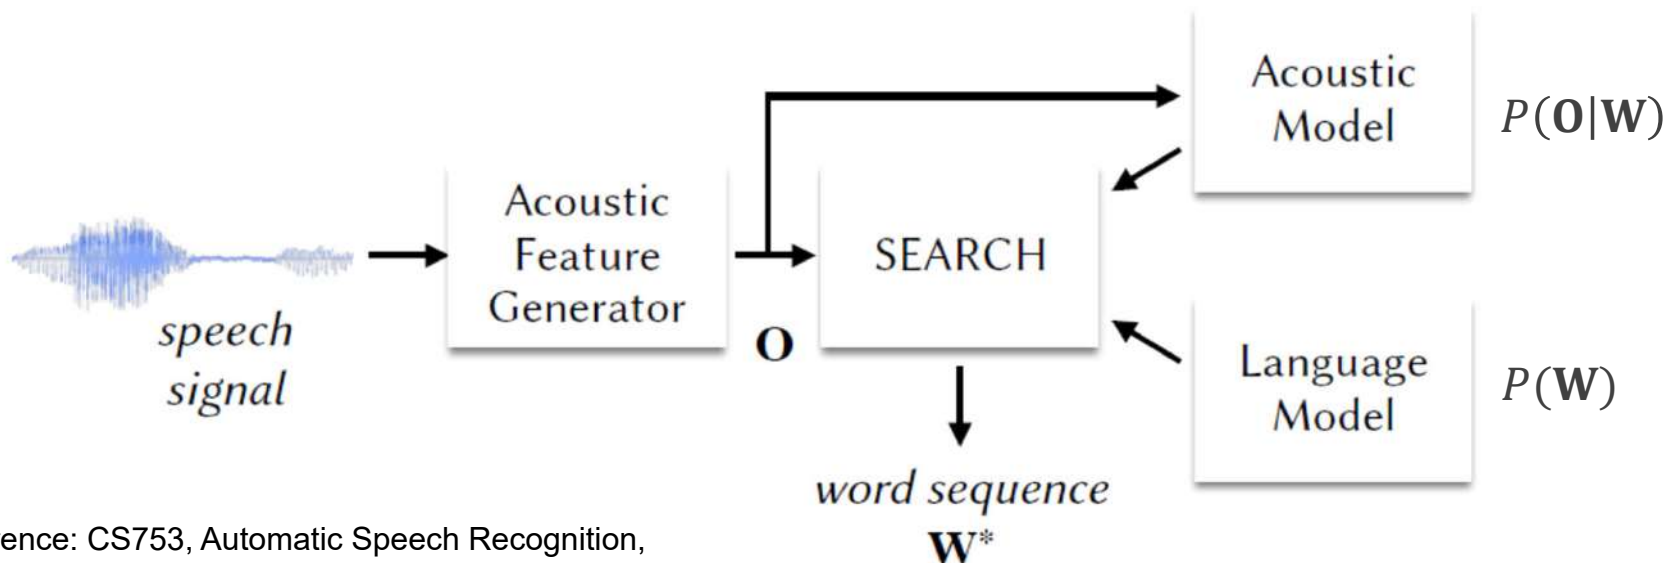




Statistical speech recognition

- Let \mathbf{O} represent a sequence of acoustic feature observations (i.e., $\mathbf{O} = \{o_1, o_2, \dots, o_t\}$), and \mathbf{W} denote a word sequence. Then the speech recognizer decodes \mathbf{W}^* as

$$\mathbf{w}^* = \underset{\mathbf{W}}{\operatorname{argmax}} P(\mathbf{W}|\mathbf{O}) = \underset{\mathbf{W}}{\operatorname{argmax}} \frac{P(\mathbf{O}|\mathbf{W})P(\mathbf{W})}{P(\mathbf{O})} \propto \underset{\mathbf{W}}{\operatorname{argmax}} P(\mathbf{O}|\mathbf{W})P(\mathbf{W})$$



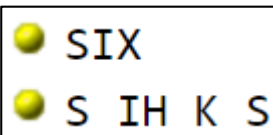
Reference: CS753, Automatic Speech Recognition,
<https://www.cse.iitb.ac.in/~pjyothi/cs753/>



Statistical speech recognition

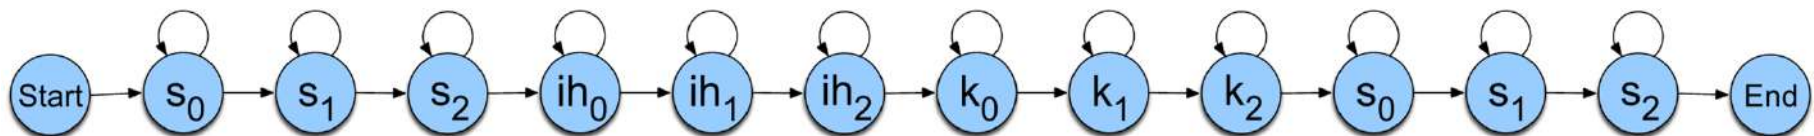
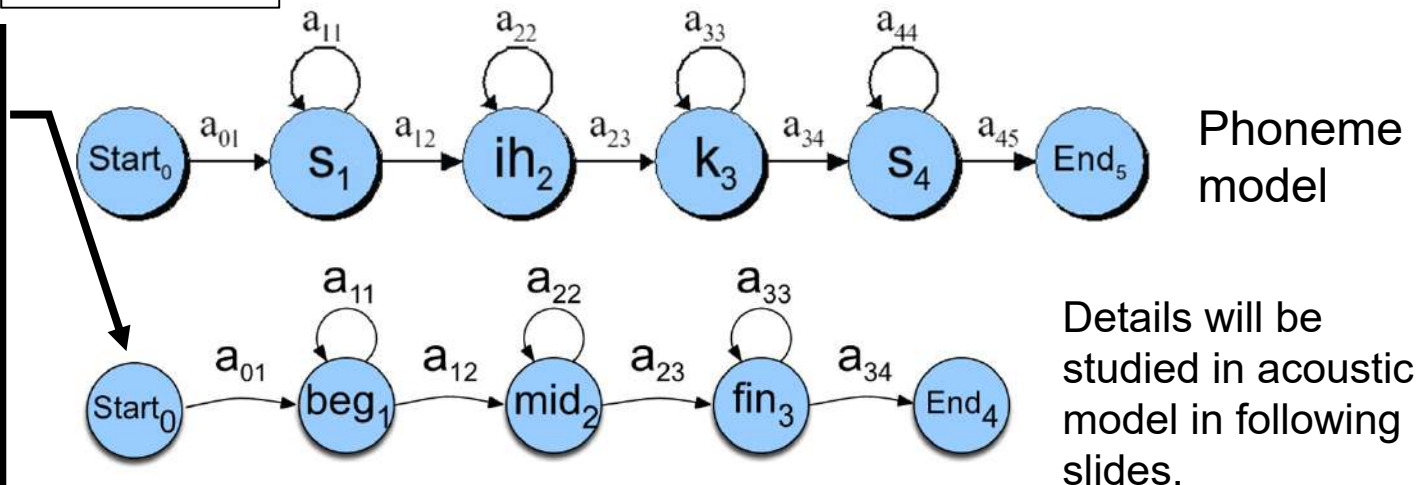
- Further introduce three definitions, acoustic signal A , phoneme L , and state Q . The optimization problem statement is changed to be

$$w^* = \underset{w}{\operatorname{argmax}} P(\mathbf{O}|\mathbf{W})P(\mathbf{W}) = \underset{w}{\operatorname{argmax}} P(\mathbf{A}|\mathbf{O})P(\mathbf{O}|\mathbf{Q})P(\mathbf{Q}|\mathbf{L})P(\mathbf{L}|\mathbf{W})P(\mathbf{W})$$



Key challenge is how to estimate sequence.
→ Hidden Markov model (HMM).

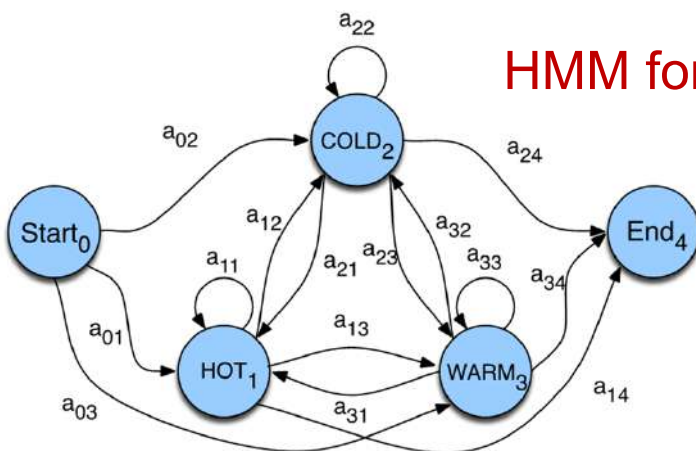
According to linguistic study, each phoneme has 3 states: (1) the transition part at the begin of the phoneme, (2) the stationary part, (3) the transition at the end.



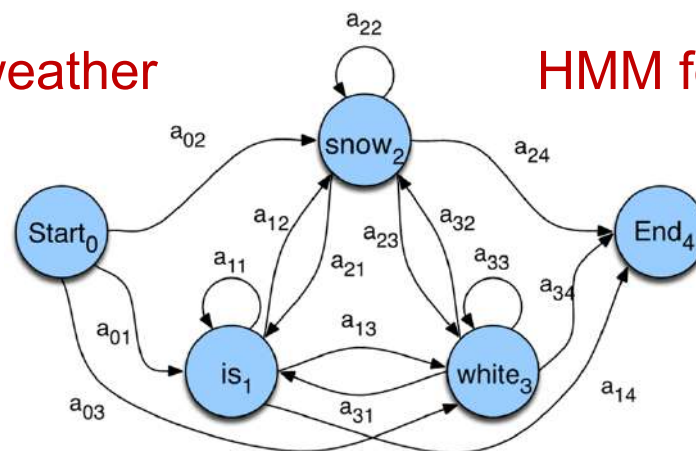
Reference: <http://www.speech.cs.cmu.edu/cgi-bin/cmudict?in=six>



Hidden Markov model (HMM): Idea



HMM for weather



HMM for word

Notation	Descriptions	
$\mathbf{Q} = \{q_1, q_2, \dots, q_t\}$	A set of N states for observations	Each observation has one state
$\mathbf{A} = \{a_{11}, a_{12}, \dots, a_{nn}\}$	A state transition probability matrix \mathbf{A} , each a_{ij} representing the probability of moving from the state i to the state j , s. t. $\sum_{j=1}^n a_{ij} = 1$	Learned from speech training dataset
$\mathbf{O} = \{o_1, o_2, \dots, o_t\}$	A sequence of T observations	Observed speech data
$B = b_i(o_t)$	An observation likelihoods , also called emission probabilities, each expressing the probability of an observation o_t being generated from a state i	Learned from speech training dataset
\mathbf{S}	A set of states (e.g., HOT ₁ , COLD ₂ , WARM ₃ , is ₁ , snow ₂ , white ₃), a special start state Start ₀ and an end state End ₄ that are not associated with observations, together with their transition probabilities out of the start state and into the end state. For example, each phoneme has 3 states (see slide 23).	

Reference: CS224S, Spoken Language Processing, <http://web.stanford.edu/class/cs224s/>

HMM: A toy example

Transition probability				Observation likelihood		
Today weather	Tomorrow weather			Weather	Probability of	
	Sunny (S)	Raining (R)	Cloudy (C)		Umbrella (U)	No umbrella (N)
Sunny (S)	0.8	0.05	0.15	Sunny (S)	0.1	0.9
Raining (R)	0.2	0.6	0.2	Raining (R)	0.8	0.2
Cloudy (C)	0.2	0.3	0.5	Cloudy (C)	0.3	0.7

Q: Given that today weather is S , what is the probability that tomorrow is S and the day after is R ?

Markov
assumption

$$P(q_2 = S, q_3 = R | q_1 = S) = P(q_3 = R | q_2 = S, q_1 = S)P(q_2 = S | q_1 = S) \\ = P(q_3 = R | q_2 = S)P(q_2 = S | q_1 = S) = 0.05 \times 0.8 = 0.04$$

Q: Given that you don't use umbrella (N) for three days, calculate the probability for the weather on these three days to be $\{q_1 = S, q_2 = C, q_3 = S\}$. Note that the prior probability for the start state as sunny (S) on day one is assumed to be $1/3$ (three weather has same probability).

$$P(q_1 = S, q_2 = C, q_3 = S | o_1 = N, o_2 = N, o_3 = N) \\ = P(o_1 = N | q_1 = S)P(o_2 = N | q_2 = C)P(o_3 = N | q_3 = S)P(q_1 = S)P(q_2 = C | q_1 = S)P(q_3 = S | q_2 = C) \\ = 0.9 \times 0.7 \times 0.9 \times 1/3 \times 0.15 \times 0.2 = 0.0057$$

Reference: <http://www.iitg.ac.in/samudravijaya/tutorials/hmmTutorialBarbaraExercises.pdf>



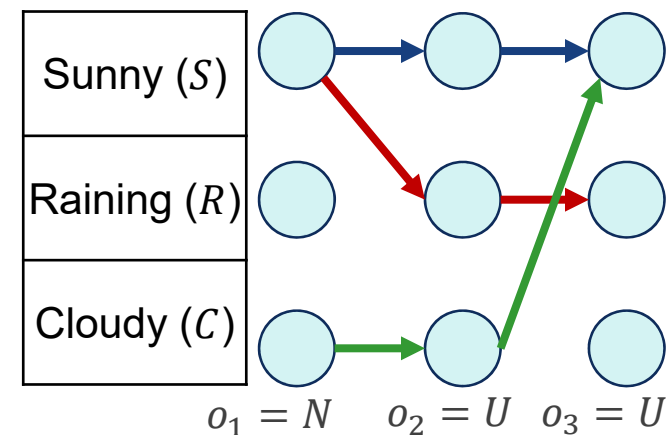
HMM: Sequence estimation

Q: Given that three days your umbrella observations are: {no umbrella (N), umbrella (U), umbrella (U)}, find the most probable weather-sequence.

Idea 1: If we ignore the weather as a 'sequence' and treat each day weather separately, the most probable weather are Sunny (S), Raining (R), Raining (R).

Idea 2: Exhaustively evaluate probability of each sequence. Consider following three possible sequences, which is most probable?

- Blue sequence: Sunny (S), Sunny (S), Sunny (S)
- Red sequence: Sunny (S), Raining (R), Raining (R)
- Green sequence: Cloudy (C), Cloudy (C), Sunny (S)



Idea 3: Design an efficient method to evaluate all possible sequence and find the most probable one.

→ We will study **Viterbi algorithm** in next few slides.

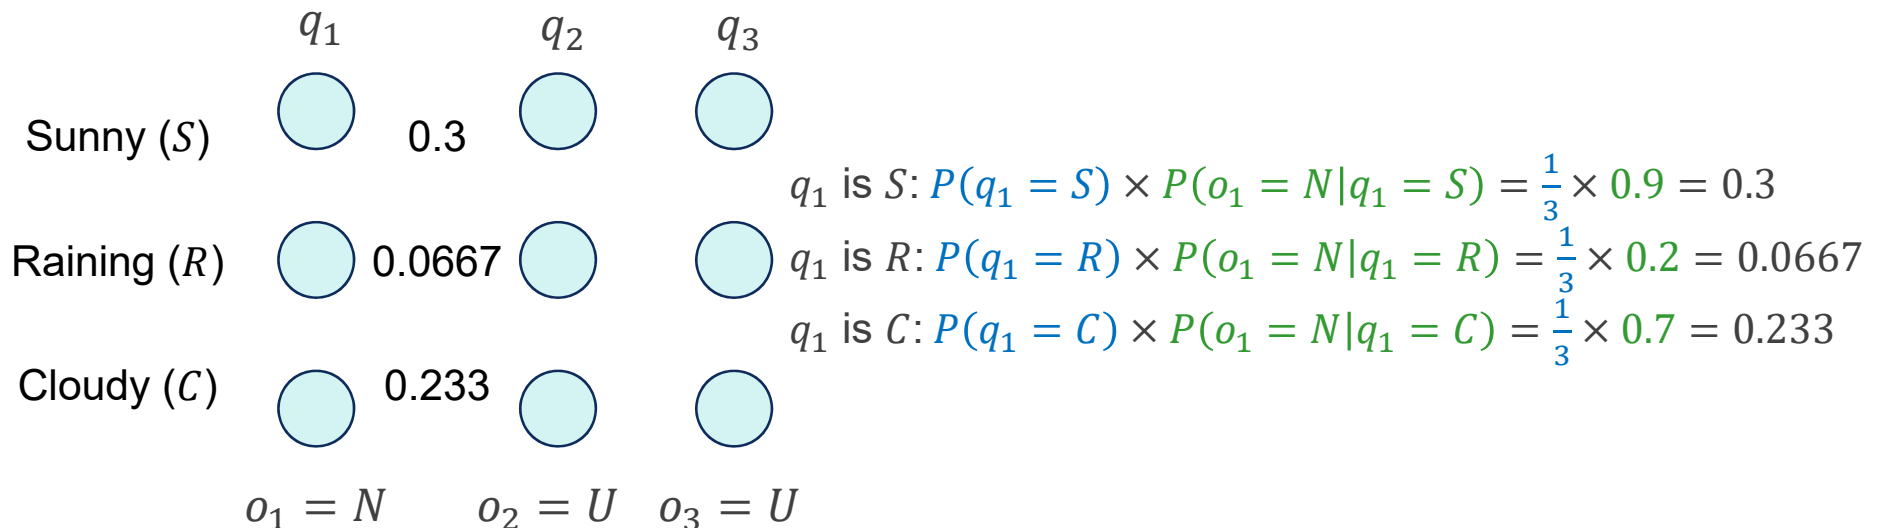
Viterbi: A single-line
.predict(0) function
in hmmlearn library



HMM: Viterbi algorithm

Key idea: “Optimal policy is composed of optimal sub-policies”.

1. **Initialization:** Calculate probability of the first day state based on first day observation and (equal) prior probability starting from all possible states.
2. **Recursion:** For all following days, calculate probability of each state based on current observation and the largest transition probability from the previous day. Record the ‘best path’ ending at current state from the previous day.
3. **Termination and back tracing:** For the last day, choose the state with the highest probability. Trace back according to the recorded most probable path.



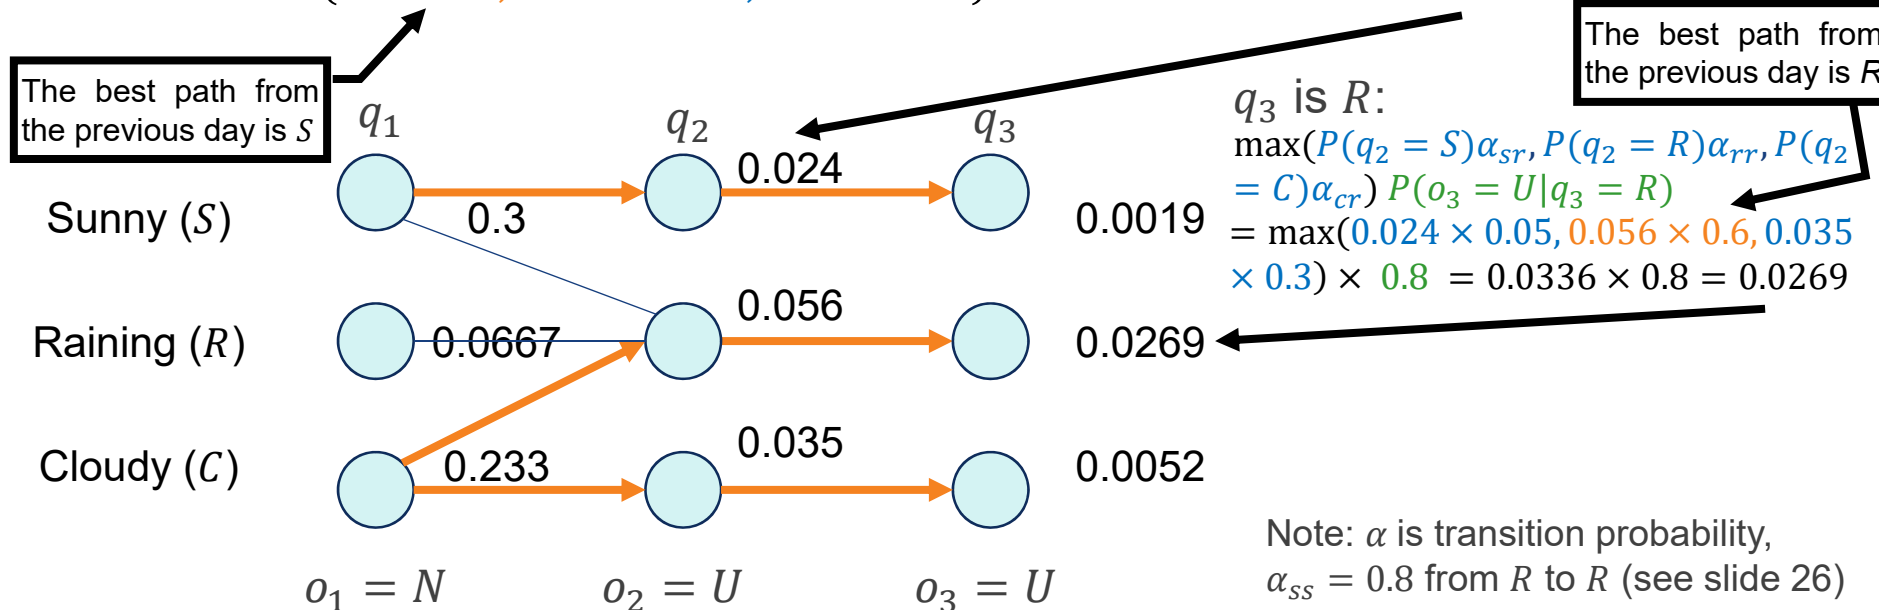


HMM: Viterbi algorithm

Key idea: “Optimal policy is composed of optimal sub-policies”.

1. Initialization: Calculate probability of the first day state based on first day observation and (equal) prior probability starting from all possible states.
2. Recursion: For all following days, calculate probability of each state based on current observation and the largest transition probability from the previous day. Record the ‘best path’ ending at current state from the previous day.
3. Termination and back tracing: For the last day, choose the state with the highest probability. Trace back according to the recorded most probable path.

$$q_2 \text{ is } S: \max(P(q_1 = S)\alpha_{ss}, P(q_1 = R)\alpha_{rs}, P(q_1 = C)\alpha_{cs}) P(o_2 = U|q_2 = S) \\ = \max(0.3 \times 0.8, 0.0667 \times 0.2, 0.233 \times 0.2) \times 0.1 = 0.24 \times 0.1 = 0.024$$



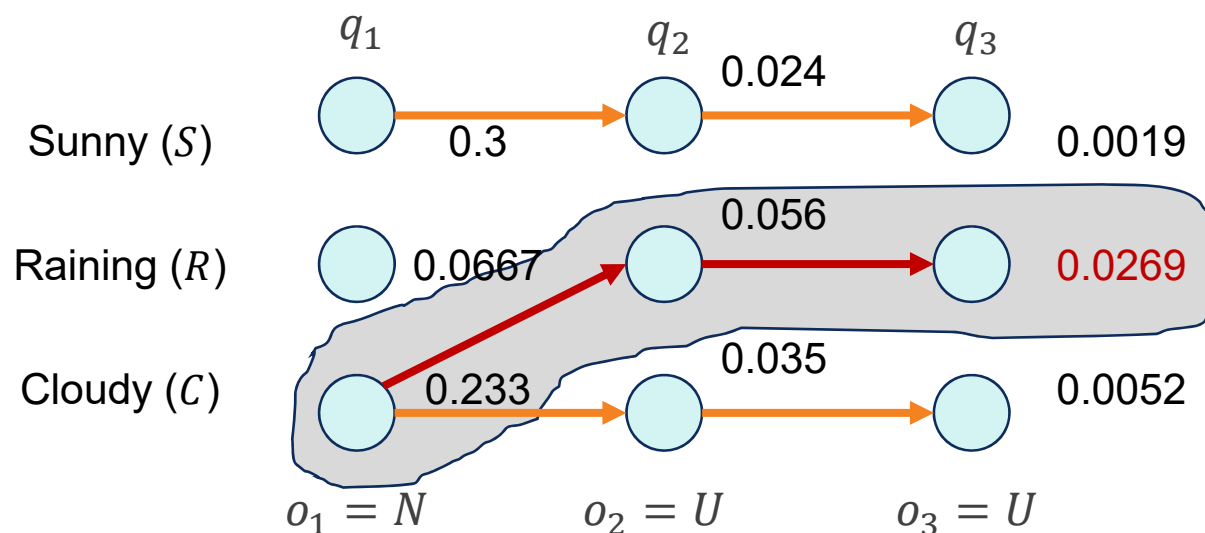
HMM: Viterbi algorithm

Key idea: “Optimal policy is composed of optimal sub-policies”.

1. Initialization: Calculate probability of the first day state based on first day observation and (equal) prior probability starting from all possible states.
2. Recursion: For all following days, calculate probability of each state based on current observation and the largest transition probability from the previous day. Record the ‘best path’ ending at current state from the previous day.
3. Termination and back tracing: For the last day, choose the state with the highest probability. Trace back according to the recorded most probable path.

The optimal sequence: Cloudy (C), Raining (R), Raining (R).

Recall that the result (in Idea 2 in slide 26) is Sunny (S), Raining (R), Raining (R).



How to use it for speech recognition?

- Weather \rightarrow phoneme state
- Umbrella \rightarrow audio features



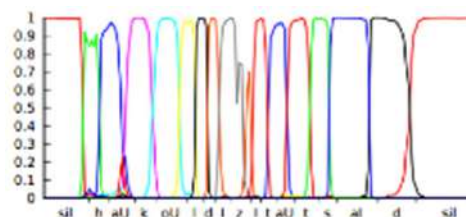
Summary: Statistical speech recognition

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmax}} P(\mathbf{O}|\mathbf{W})P(\mathbf{W}) = \underset{\mathbf{w}}{\operatorname{argmax}} P(\mathbf{A}|\mathbf{O})P(\mathbf{O}|\mathbf{Q})P(\mathbf{Q}|\mathbf{L})P(\mathbf{L}|\mathbf{W})P(\mathbf{W})$$



A

Speech Audio

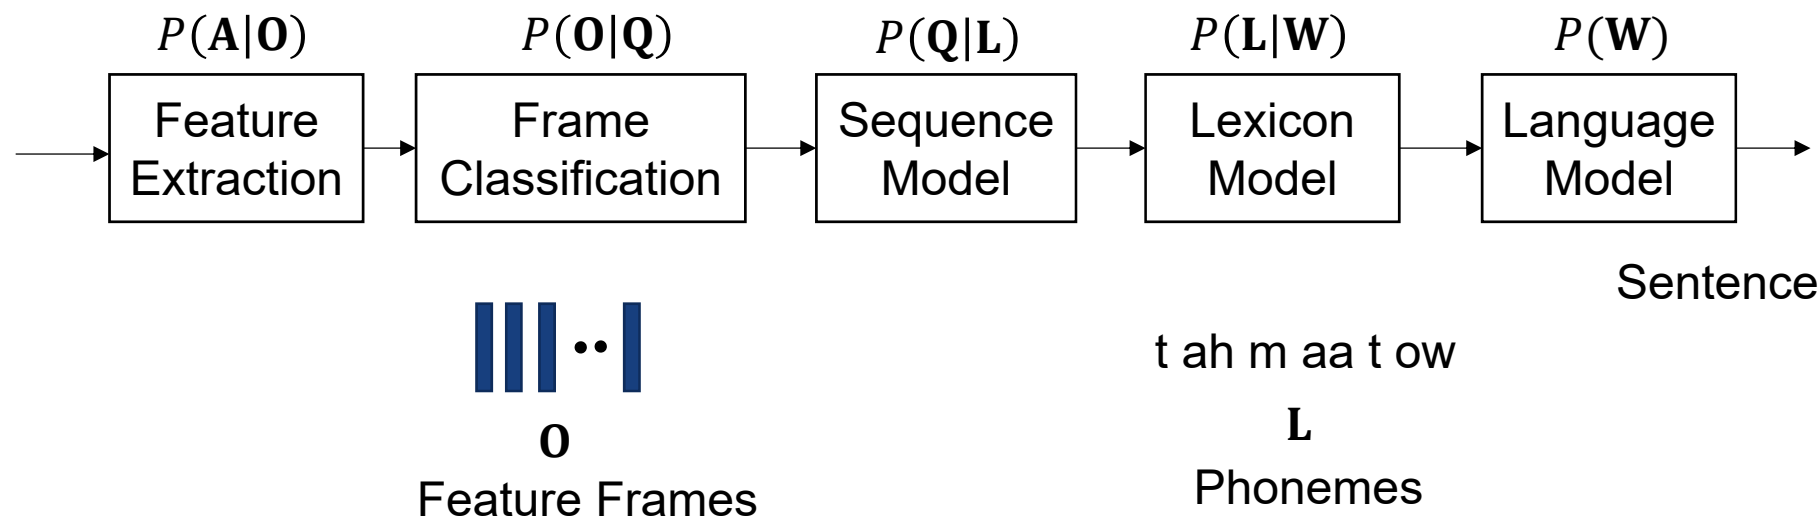


Q

Sequence States

W

Words





Statistical speech recognition: Lexicon model

- Lexical modelling forms the bridge between the acoustic and language models
- Each one with a pronunciation in terms of phones
- CMU dictionary: 127K words,
<http://www.speech.cs.cmu.edu/cgi-bin/cmudict>

Deterministic model

Word	Pronunciation
TOMATO	t ah m aa t ow
	t ah m ey t ow
COVERA GE	k ah v er ah jh
	k ah v r ah jh

Probabilistic model

Word	Pronunciation	Probability
TOMATO	t ah m aa t ow	0.45
	t ah m ey t ow	0.55
COVERA GE	k ah v er ah jh	0.65
	k ah v r ah jh	0.35

Statistical speech recognition: Language model

N-gram models: Build the language model by calculating probabilities from text training corpus: how likely is one word to follow another.

	i	want	to	eat	chinese	food	lunch	spend
i	5	827	0	9	0	0	0	2
want	2	0	608	1	6	6	5	1
to	2	0	4	686	2	0	6	211
eat	0	0	2	0	16	2	42	0
chinese	1	0	0	0	0	82	1	0
food	15	0	15	0	1	4	0	0
lunch	2	0	0	0	0	1	0	0
spend	1	0	1	0	0	0	0	0

Example: Bi-gram counts for eight of the words (out of $V = 1446$) in the Berkeley Restaurant Project corpus of 9332 sentences

can you tell me about any good cantonese restaurants close by
mid priced thai food is what i'm looking for
tell me about chez panisse
can you give me a listing of the kinds of food that are available
i'm looking for a good place to eat breakfast
when is caffe venezia open during the day



Statistical speech recognition: Data

- Collect corpora appropriate for recognition task
 - Small speech + phonetic transcription to associate sounds with symbols (Acoustic Model)
 - Large (≥ 60 hrs) speech + orthographic transcription to associate words with sounds (Acoustic Model)
 - Very large text corpus to identify N -gram probabilities or build a grammar (Language Model)



Speech recognition: Evaluation

- **Word Error Rate** (WER): Minimum Edit Distance: Distance in words between the system output and the reference transcription (truth)

$$WER = \frac{S + D + I}{N}$$

- S is the number of substitutions,
- D is the number of deletions,
- I is the number of insertions and
- N is the number of words in the reference

Truth: What a bright day

System: What a day

Deletion: "Bright" was deleted by the system

Truth: What a day

System: What a bright day

Insertion: "Bright" was inserted by the system

Truth: What a bright day

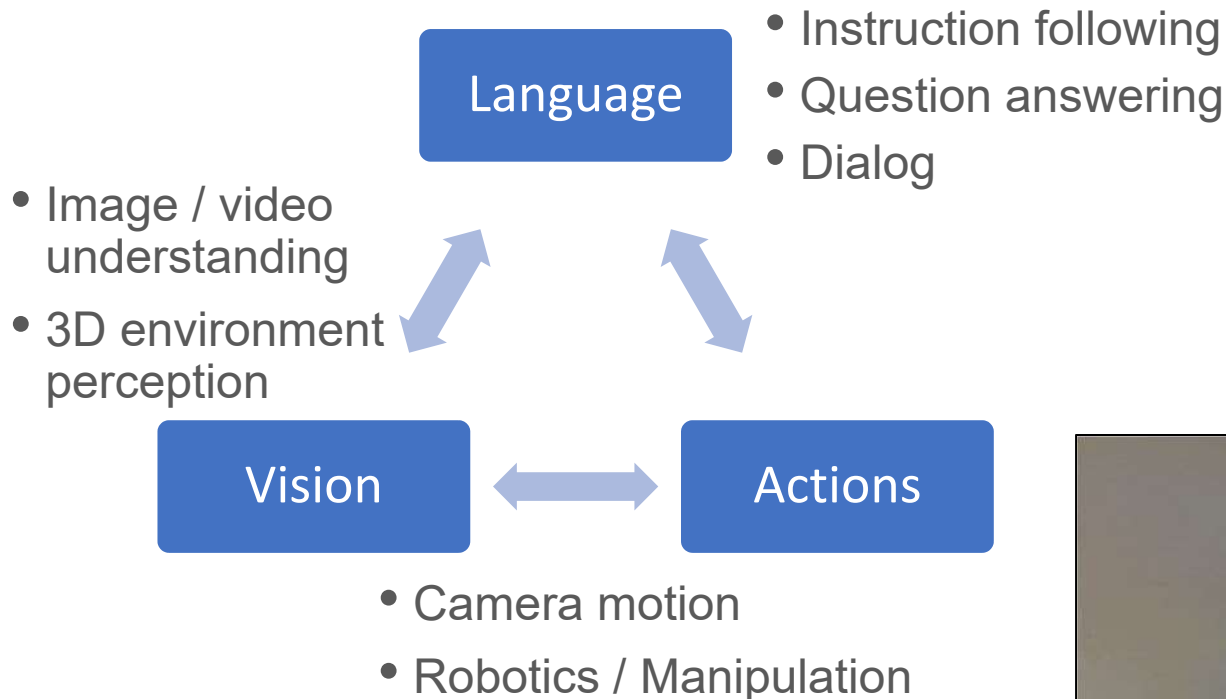
System: What a light day

Substitution: "Bright" was substituted by "light" by the system

Reference: <https://martin-thoma.com/word-error-rate-calculation/>



Language, vision and actions



Reference:

- Connecting language and vision to actions, <https://lvatutorial.github.io/>
- Reference: Natural language interaction with robots, <https://bringmeaspoon.org/>



Visual question answering (VQA)

- **Task:** Given an image and a natural language open-ended question, generate a natural language answer. This is reasoning techniques using both language and vision knowledge.

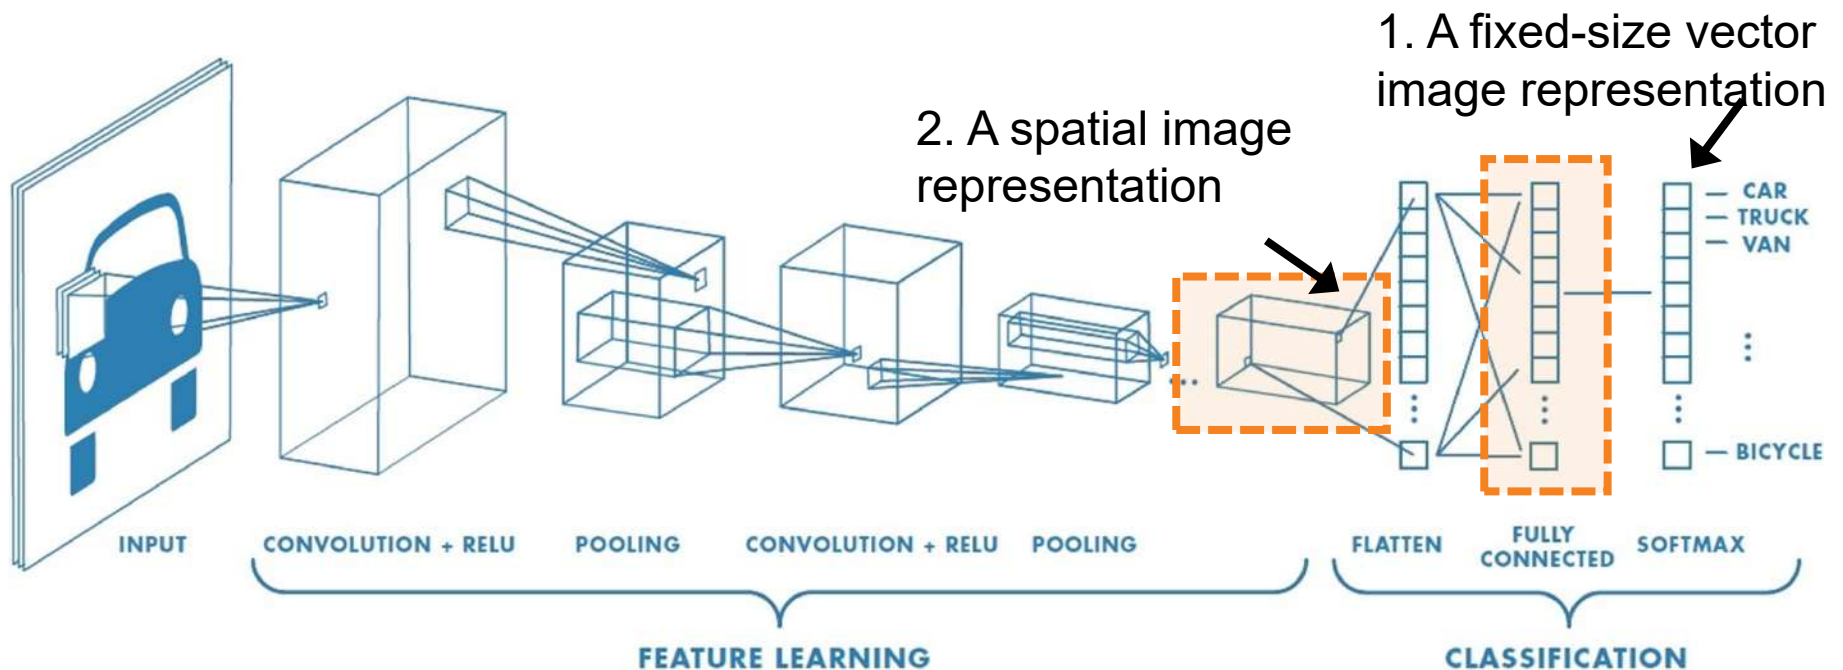


Reference: P. Wang, Q. Wu, C. Shen, A. Hengel, A. Dick, FVQA: Fact-Based Visual Question Answering, <https://arxiv.org/abs/1606.05433>



Visual question answering (VQA)

Recap: We can treat convolutional neural networks (CNNs) as black boxes that can output following two things



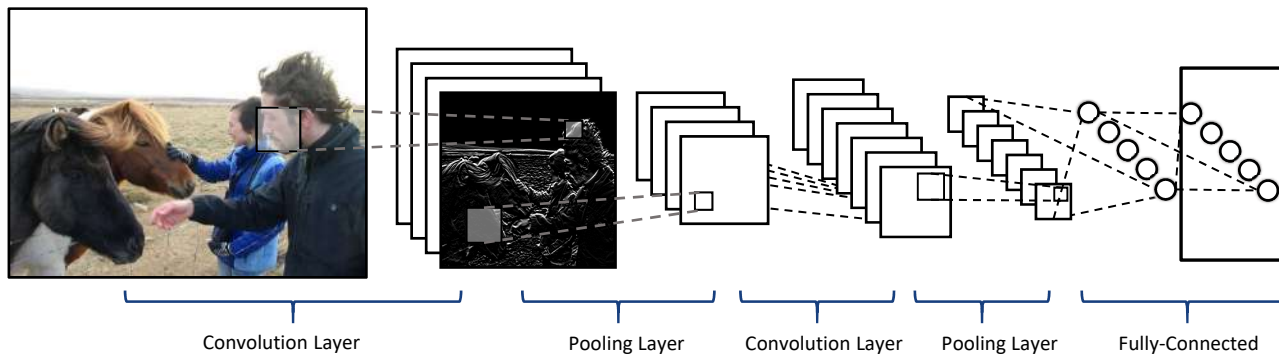
Reference: Connecting language and vision to actions, <https://lvatutorial.github.io/>



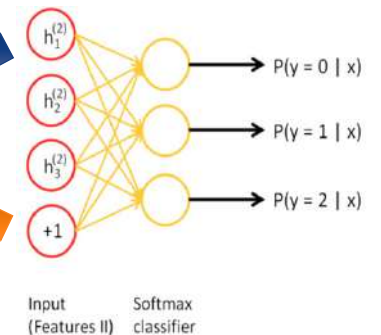
Visual question answering (VQA)

Demo website: <http://vqa.cloudcv.org/>

Image

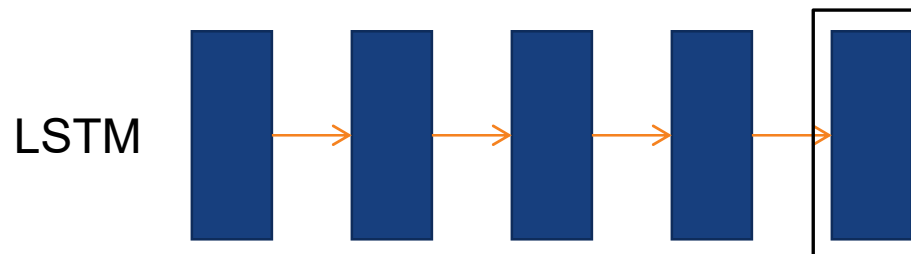


Neural Network over top answers



Question

"How many horses are in this image?"



Reference: Aishwarya Agrawal, et al, VQA: Visual Question Answering, <https://arxiv.org/abs/1505.00468>



Visual question answering (VQA)

- Output: (classification) 1000 answers.
- Input: Image feature (VGG16 model, 4096 dimensional vector, 'fc2'), text feature (tokens in the question are first embedded into 300 dimensional GloVe vectors and then passed through LSTM). Both multimodal data points are then passed through a dense layer and a final softmax layer.

Layer (type)	Output Shape	Param #
reshape_2_input (InputLayer)	(None, 4096)	0
input_2 (InputLayer)	(None, 30, 300)	0
reshape_2 (Reshape)	(None, 4096)	0
lstm_2 (LSTM)	(None, 512)	1665024
dense_4 (Dense)	(None, 1024)	4195328
concatenate_2 (Concatenate)	(None, 1536)	0
dense_5 (Dense)	(None, 1024)	1573888
dense_6 (Dense)	(None, 1000)	1025000
Total params: 8,459,240		
Trainable params: 8,459,240		
Non-trainable params: 0		

Toy model architecture



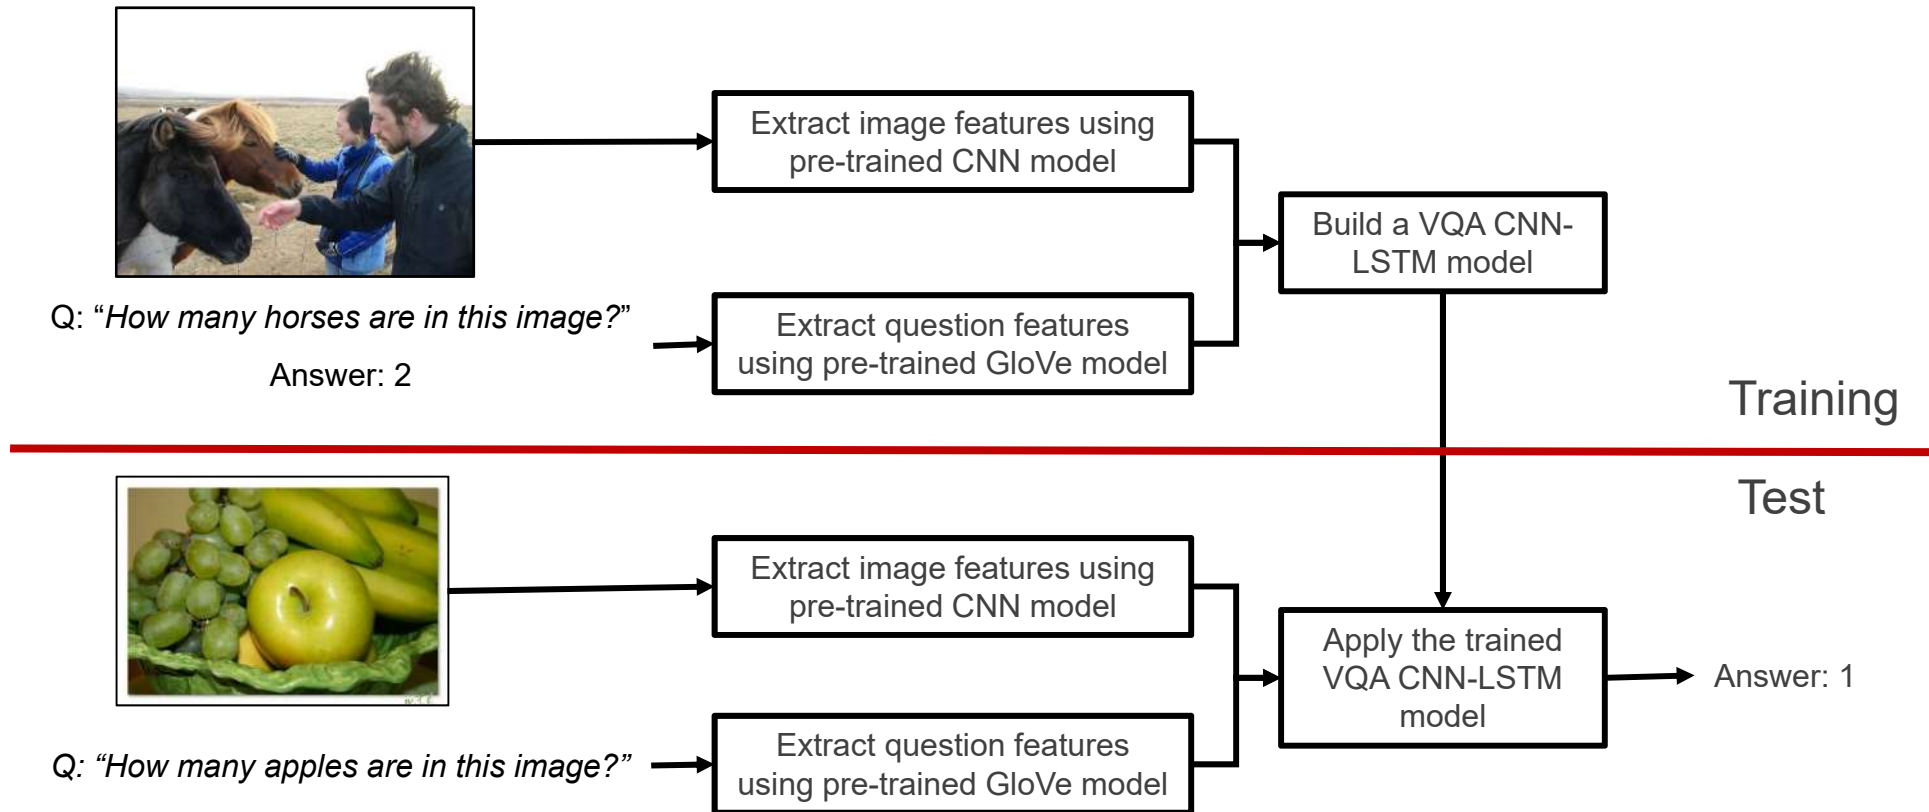
Example training images and question/answer



Visual question answering (VQA)

The full VQA pipeline

- Pre-process both image and question/answer text
- Design a model architecture and train the model
- Deploy the model and process the new test image and question input





Workshop: Speech cognitive systems

Objective

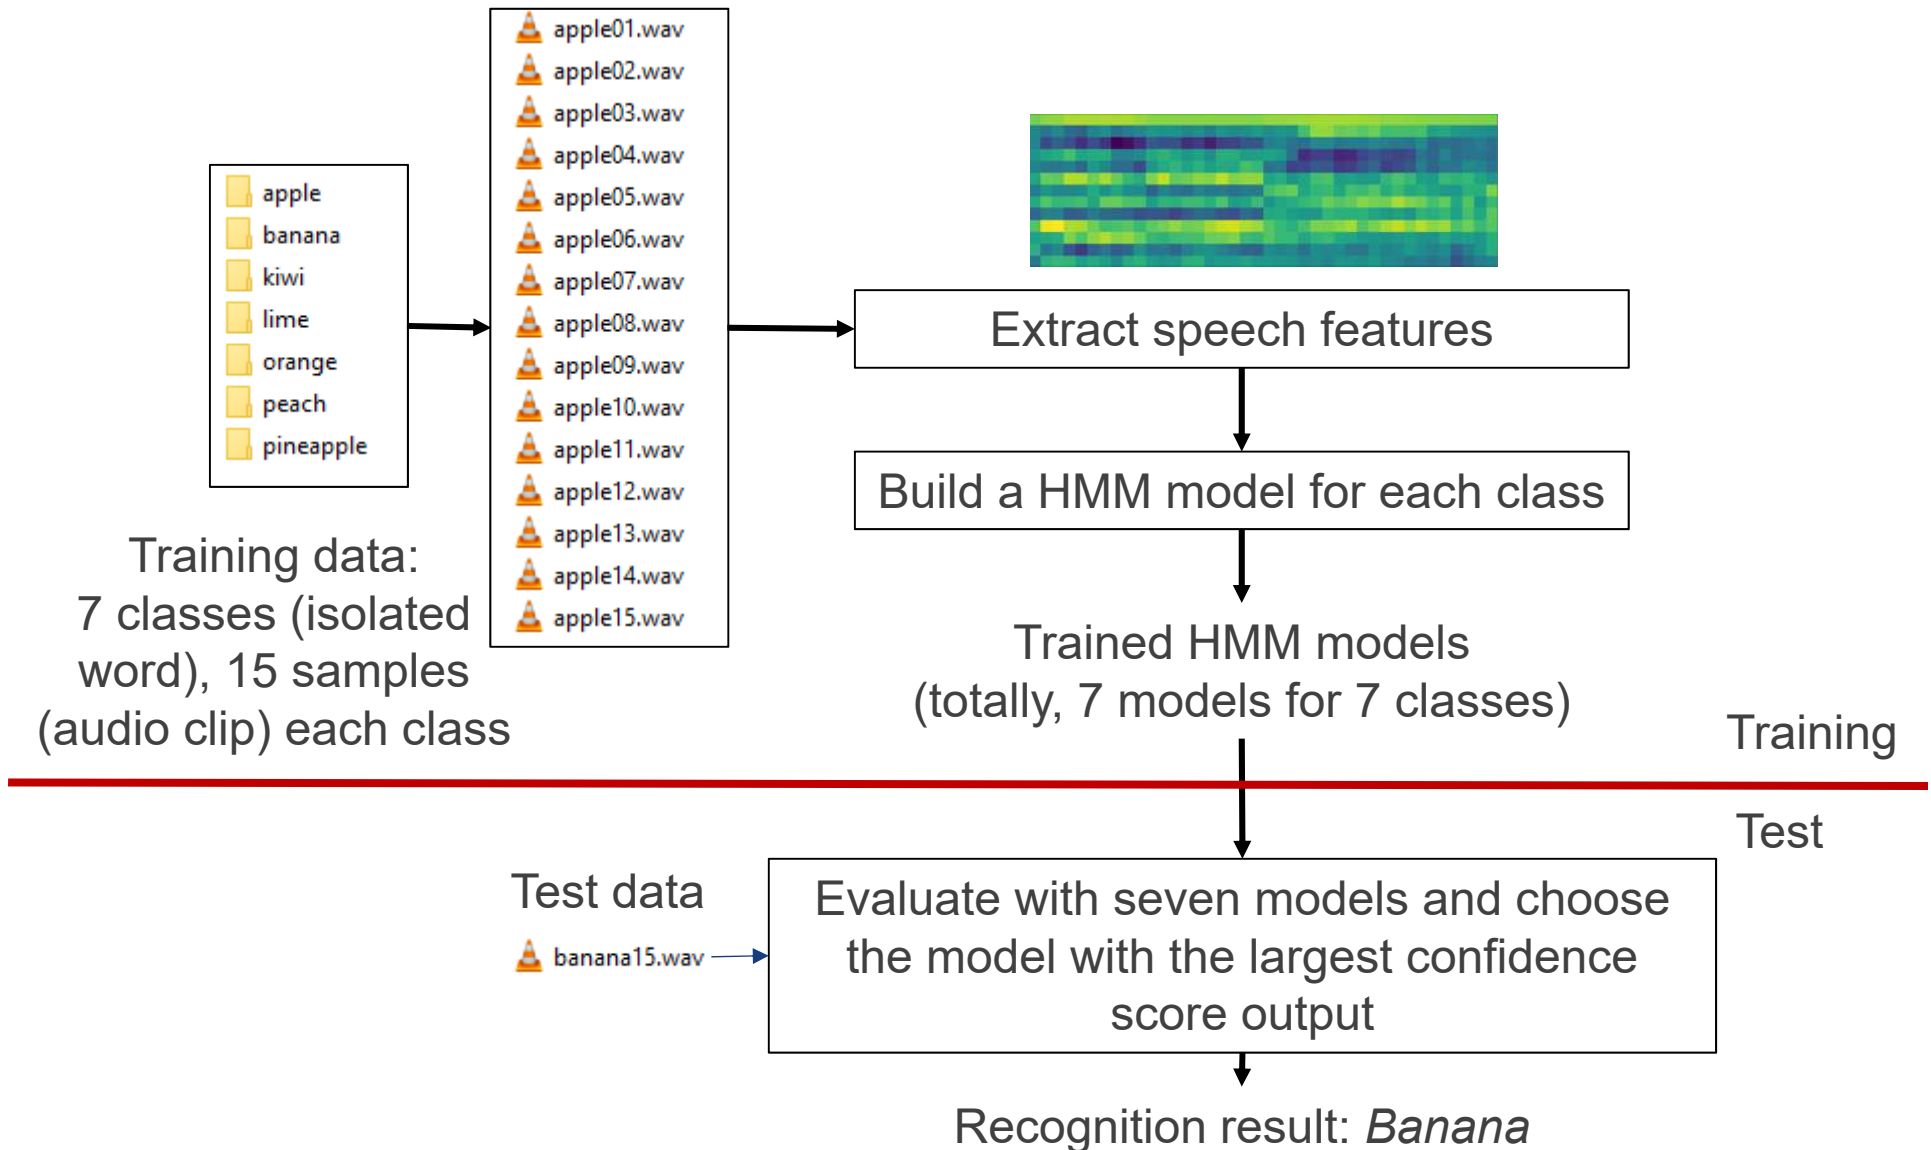
- Build a HMM-based speech recognizer

Reference

- Prateek Joshi, Python Machine Learning Cookbook, Packt Publishing, 2016, Code available at <https://github.com/PacktPublishing/Python-Machine-Learning-Cookbook>

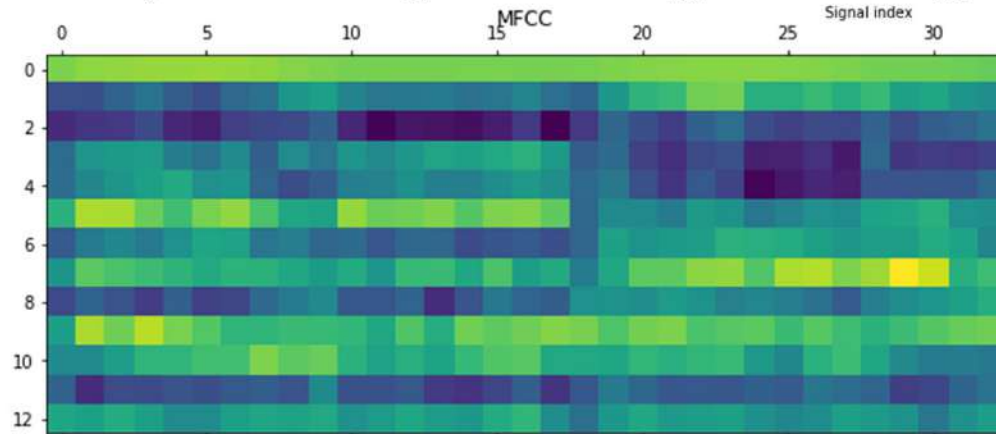
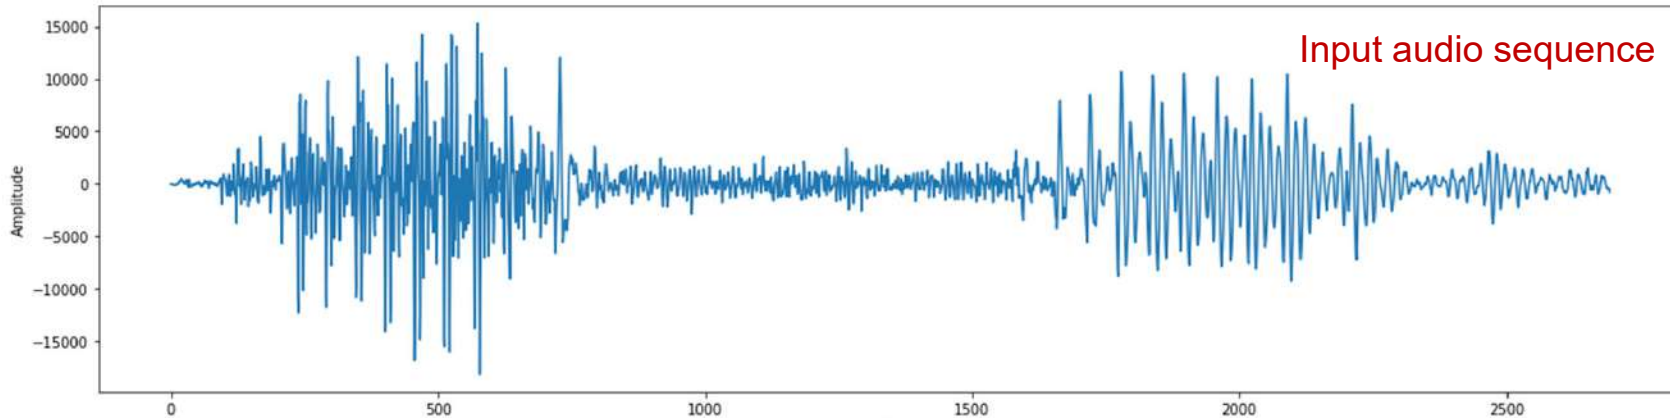


Workshop: Speech cognitive systems





Workshop: Speech cognitive systems



Variables	Dimensions
Input audio (apple01.wav)	(2694,)
MFCC feature of input signal (# windows, # features)	(33, 13)
HMM transition probability matrix with 3 components	(3, 3)
HMM state sequence	(33,)

```
[[9.32529578e-01, 4.36394206e-26, 6.74704225e-02]
 [1.53171049e-39, 9.05902521e-01, 9.40974787e-02]
 [1.24927288e-01, 1.02681348e-01, 7.72391365e-01]]
```

HMM transition probability matrix

```
[0 0 0 0 0 0 0 0 0 2 0 0 0 0 0 0 0 2 2 1 1 1 1 1 1 1 1 1 1 1 2 2]
```

HMM state sequence

DEMO: Colab

- Unzip the demo file in your local machine.
- Create a new **folder** (say, AIOTS) in your Google drive.
- Upload your unzipped local demo files into Google drive.
- Select the .ipynb file, Right click, Open with Google Colaboratory

The screenshot illustrates the process of opening a Jupyter Notebook file from Google Drive using Google Colaboratory. The Google Drive interface shows a file named 'wk_speech_recognition_Colab_v1.0.ipynb' in the 'CGS' folder. A right-click context menu is open, and the 'Open with' option is selected, leading to a submenu where 'Google Colaboratory' is highlighted. Another submenu shows '+ Code' and '+ Text' options. Below this, the Jupyter Notebook interface is shown, displaying the installation of packages and the execution of code to mount Google Drive and change the working directory.

```
Installing collected packages: python-speech-features, hmmlearn
Successfully installed hmmlearn-0.2.1 python-speech-features-0.6

1 # Mount your drive
2 # Run this cell, then you'll see a link, click on that link, allow access
3 # Copy the code that pops up, Paste it in the box, Hit enter
4
5 from google.colab import drive
6 drive.mount('/content/gdrive')
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

```
[6] 1 # Change working directory to be current folder
    2 import os
    3 os.chdir('/content/gdrive/My Drive/CGS')
    4 !ls

data wk_speech_recognition_Colab_v1.0.ipynb
```



What we have learnt

- A typical vision cognitive system pipeline
- A statistical speech cognitive system framework
- Isolated word speech recognition using Hidden Markov model (HMM)

Thank you!

Dr TIAN Jing
Email: tianjing@nus.edu.sg